

NOTE Run all cells in section 2, then skip to sections 5 for data loading (all preprocessing is done) and modeling helpers. Then, be sure to run the first cell in each other section to ensure the data is formatted correctly for that model type



pneumonia image

Business Problem

Pneumonia is a common infection that causes inflammation and possible fluid accumulation in the air sacs of the lungs. In China, pneumonia is one of the leading causes of death for [children under 5 years old](#). The infection is commonly caused by bacteria and viruses, but can also be caused by fungal sources.

Pediatric pneumonia is initially diagnosed based on the [time of the year and the results of a physical exam](#), paying attention to the child's breathing and listening to the lungs. Physical symptoms associated with pneumonia generally include fever, rapid breathing, and increased heart rate. A further step in diagnosis would be to use chest X-rays; pneumonia is not always seen on x-rays, either because [the disease is in its initial stages or it involves a part of the lung not easily seen by X-ray](#). Inconclusive initial testing can result in additional blood tests to confirm or rule out the presence of infection.

As we can see, even with modern medicine pneumonia can be misdiagnosed. Not all pneumonia infections are treated in the same way; viral infections cannot be fought with the same antibiotics that would treat a bacterial infection. A fast and accurate diagnosis allows doctors to treat the infection with the appropriate care. One application of machine learning in medicine is digital diagnosis. We have been tasked with developing an identification model to determine if a chest X-ray indicates the presence of pneumonia. False negative results are to be minimized compared to false positives.

The data is sourced from [Kaggle](#). It is already split into three folders for training, validation and testing. All the chest radiographs were screened for quality and diagnostic labeling performed by physicians. The images were collected during routine clinical care of pediatric patients between one and five years old from Guangzhou Women and Children's Medical Center in Guangzhou, China.

Imports

```
In [1]: import os
        from os import listdir
        from os.path import isfile, join
        import tempfile

        from PIL import Image
        import math

        import pandas as pd
        import numpy as np
```

```

import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
%matplotlib inline

import sklearn
from sklearn.model_selection import cross_val_score, StratifiedKFold
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.utils.class_weight import compute_class_weight

import tensorflow as tf
from tensorflow import keras

from keras.callbacks import EarlyStopping, ModelCheckpoint
from keras.wrappers.scikit_learn import KerasClassifier
from keras.models import Sequential
from keras.layers import Dense, Dropout, Conv2D, MaxPooling2D, Flatten

```

In [2]:

```

# set figsize for matplotlib
mpl.rcParams['figure.figsize'] = (12, 10)
# set colors for plots
colors = plt.rcParams['axes.prop_cycle'].by_key()['color']

```

EDA

Helper Functions

Helper functions:

- get_ratio: return the ratio of an 2D image array width:height
- resize_and_crop: returns a scaled and cropped image

In [3]:

```

def get_ratio(array):
    return array.shape[0]/array.shape[1]

def resize_and_crop(image, new_size=256):
    """
    This function will take in an image, resize and crop it to a square.
    The default new length for a side is 256 pixels.
    The function will output a new square, centered image.

    """
    # get width and height from passed image
    width, height = image.size
    # set up variables for new width and height
    new_width = 0
    new_height = 0
    # calculate the ratio and length of the smaller dimension if the larger moves to th
    if width > height:
        ratio_wh = width / height
        new_width = int(ratio_wh * new_size)
        new_height = new_size
    else:
        ratio_hw = height / width

```

```

        new_width = new_size
        new_height = int(ratio_hw * new_size)
    # resize the image
    scaled_image = image.resize((new_width, new_height))
    # if the image is wider, crop in equally from the sides to preserve center of image
    # perform the opposite on images that are taller
    if new_width > new_height:
        # we are only cropping in from left and right, so set top crop to 0 and bottom
        top = 0
        bottom = new_size
        # set the left and right side crop values
        left = int(math.ceil((new_width - new_size) / 2))
        right = new_width - int(math.floor((new_width - new_size) / 2))
        # crop the image
        cropped_image = scaled_image.crop((left, top, right, bottom))
    else:
        # we are only cropping top and bottom, so set left crop to 0 and right crop to
        left = 0
        right = new_size
        # set the top and bottom crop values
        top = int(math.ceil((new_height - new_size) / 2))
        bottom = new_height - int(math.floor((new_height - new_size) / 2))
        # save the cropped image
        cropped_image = scaled_image.crop((left, top, right, bottom))
    # return the scaled and cropped image
    return cropped_image

```

Image size

First, we are going to build a dataframe that contains all the image information for us to look at image sizes.

```

In [4]: folder_names = ['train', 'test', 'val']
        label_names = ['NORMAL', 'PNEUMONIA']

        all_images = []

        # Loop through the different combinations of folder name prefixes
        for folder in folder_names:
            for label in label_names:
                # set up the path to each folder of images
                path = f'./chest_xray/{folder}/{label}'
                # create a list of the filenames in that directory
                filelist = list(listdir(path))
                # Loop through each file in the folder
                for file_name in filelist:
                    # set the filepath for the file in question
                    filepath = path + r'/' + file_name
                    # open the image
                    image = Image.open(filepath)
                    # save the image mode
                    img_mode = image.mode
                    # get the width and height of the image
                    width, height = image.size
                    # calculate the dimension ratio
                    ratio = height / width
                    # append the image information to the list all_images

```

```

all_images.append((folder, label, file_name, img_mode, ratio, width, height)

# set the column names
column_names = ['folder', 'label', 'file_name', 'img_mode', 'ratio', 'width', 'height']
# convert the list of images to a dataframe
all_images_df = pd.DataFrame(all_images, columns=column_names)

```

In [5]: `all_images_df.head()`

Out[5]:

	folder	label	file_name	img_mode	ratio	width	height
0	train	NORMAL	IM-0115-0001.jpeg	L	0.888995	2090	1858
1	train	NORMAL	IM-0117-0001.jpeg	L	0.810127	1422	1152
2	train	NORMAL	IM-0119-0001.jpeg	L	0.792265	1810	1434
3	train	NORMAL	IM-0122-0001.jpeg	L	0.790482	1618	1279
4	train	NORMAL	IM-0125-0001.jpeg	L	0.703125	1600	1125

In [6]: `all_images_df.img_mode.value_counts()`

Out[6]:

```

L      5573
RGB     283
Name: img_mode, dtype: int64

```

So it looks like despite our images being in black and white, we have 283 images that are coded as RGB images. We will need to convert them before making our image arrays otherwise our shapes will not match. Lets look at the sizes of our images.

In [7]: `all_images_df.describe()`

Out[7]:

	ratio	width	height
count	5856.000000	5856.000000	5856.000000
mean	0.712905	1327.880806	970.689037
std	0.117312	363.500922	383.392117
min	0.295964	384.000000	127.000000
25%	0.630616	1056.000000	688.000000
50%	0.706272	1281.000000	888.000000
75%	0.792627	1560.000000	1187.000000
max	1.197044	2916.000000	2713.000000

We have a wide variety of image sizes. For best results in modeling it was advised to scale images to 256x256. We can see that there are some images whose height is below our minimum, lets look a little closer.

In [8]: `small_images_df = all_images_df[all_images_df.height < 256].copy()`

```
In [9]: small_image_count = len(small_images_df)

print(f'There are {small_image_count} images with a dimension under 256.')
print('-----')
print('Small images occur in the following data/labels:')
display(all_images_df[all_images_df.height < 256].folder.value_counts())
display(all_images_df[all_images_df.height < 256].label.value_counts())
```

There are 58 images with a dimension under 256.

Small images occur in the following data/labels:

```
train      58
Name: folder, dtype: int64
PNEUMONIA   58
Name: label, dtype: int64
```

```
In [10]: print('Target distribution in training data')
round(all_images_df[all_images_df.folder == 'train'].label.value_counts(normalize=True))
```


Target distribution in training data

```
Out[10]: PNEUMONIA    74.29
NORMAL      25.71
Name: label, dtype: float64
```

```
In [11]: percent_small = round(small_image_count / len(all_images_df[(all_images_df.label == "PN
print(f'The {small_image_count} small images represent {percent_small}% of our training
```

The 58 small images represent 0.01% of our training data with the pneumonia label

```
In [12]: small_images_df.ratio.plot(kind='hist');
```



We can see here that there are only 58 images out of our total 5856 which have a dimension less than our minimum. They are all from training images positive with pneumonia, where we have a large imbalance of that data. Additionally, the ratios of these images are small, meaning that they are significantly wider than they are tall. This would also pose a problem as our images are chest x-rays, and those wider images would likely get cut off when we resize them to square.

We will not consider these images in modeling so we will drop them from the dataframe.

```
In [13]: images_df = all_images_df.drop(small_images_df.index)
images_df.reset_index(drop=True, inplace=True)
```


Since we discovered that the smaller images had very low ratios, let's consider the ratios of the all the rest of the images.

```
In [14]: images_df.ratio.describe()
```

```
Out[14]: count    5798.000000
mean       0.715844
std        0.113936
min        0.387387
25%        0.632822
50%        0.707828
```

```
75%          0.793645
max          1.197044
Name: ratio, dtype: float64
```

```
In [15]: images_df.ratio.plot(kind='hist');
```



Most of our images (75%) have a ratio above 0.63, but the minimum is very low at 0.38. It's suspected that cropping the images that we have with low ratio (which indicates very wide images) may be problematic, so it may be something we have to return to and consider dropping some images with a ratio below a certain threshold. We can do some quick investigation into that.

Cropping

```
In [16]: def compare_cropping(index_value, df=images_df):
        """
        This helper function will take in the index value of an image in the default
        dataframe: all_images_df. It will display the width, height, and ratio of the
        original with the original image. It will then perform the cropping and display
        the new imaged data along with the new cropped image for comparison.
        """
        # Get the folder and img_type value from the dataframe to construct the folder root
        folder = df.iloc[index_value]['folder']
        label = df.iloc[index_value]['label']
        # construct the folder root for the file in question
        root=f"./chest_xray/{folder}/{label}/"
```

```

# set the image path and open the image
image_path = root + df.iloc[index_value]['file_name']
image = Image.open(image_path)
# print the dimensions, ratio, and original image
print(f"width: {image.width}\theight: {image.height}\tratio: {image.height/image.wi
display(image)
# crop the image
cropped_image = resize_and_crop(image)
# print the new image (will be 256x256 with ratio of 1.0)
display(cropped_image)

```



Let's look at an image with our 25th percentile value for ratio: 0.63

```

In [17]: compare_cropping(images_df[images_df.ratio == 0.63].index[0])

```

width: 1400 height: 882 ratio: 0.63



Here we can see we are losing some of our image on the sides of the chest cavity. We know that



data is expensive, so we want to use the most of our images that we can. For arguments' sake, if we were willing to lose the worst 5% of our images based on ratio, let's examine an image with the ratio of the 5% quantile.

```
In [18]: images_df.quantile(q=0.05)
```

```
Out[18]: ratio      0.54482  
width      860.85000  
height     528.00000  
Name: 0.05, dtype: float64
```

```
In [19]: compare_cropping(images_df[images_df.ratio == 0.55].index[0])
```

width: 1440 height: 792 ratio: 0.55



We are definitely losing data by cropping images with ratios this low. Out of curiosity, let's see where most of these images are.

```
In [20]: images_df[images_df.ratio <= 0.55].folder.value_counts()
```

```
Out[20]: train    283
        test     44
        Name: folder, dtype: int64
```

```
In [21]: images_df[images_df.ratio <= 0.55].label.value_counts()
```

```
Out[21]: PNEUMONIA    321
        NORMAL        6
        Name: label, dtype: int64
```

If we see impact during modeling, we may consider dropping images with ratios less than 0.55 from consideration. Most of the images are from our training set of positive cases (which we have a large majority of).

Image Preprocessing

Now that we have a dataframe of the raw images data, we need to resize, shrink and convert the images to arrays for modeling. We will do that with a helper function applied to the dataframe and input the arrays into a new feature called 'image_array'.

```
In [22]: def create_image_array(df):
        # set path for image
        path = path = f'./chest_xray/{df.folder}/{df.label}/{df.file_name}'
        # open image
        image = Image.open(path)
        # convert to greyscale
        if image.mode == 'RGB':
            image = image.convert('L')
        # resize and crop the image
        image_modified = resize_and_crop(image)
        # return the image as a numpy array
        return np.asarray(image_modified)

        images_df['image_array'] = images_df.apply(create_image_array, axis=1)
```

We also need to convert the label to a binary value for our target.

```
In [23]: def create_target(df):
        if df.label == 'PNEUMONIA':
            return 1
        else:
            return 0

        images_df['target'] = images_df.apply(create_target, axis=1)
```

Now we can set a dataframe for each of our three datasets (train, test and validation)

```
In [24]: train_df = images_df[images_df.folder == 'train'][['image_array', 'target']]
        test_df = images_df[images_df.folder == 'test'][['image_array', 'target']]
        val_df = images_df[images_df.folder == 'val'][['image_array', 'target']]
```

Then we split the image_array and target apart, save each datasets' X and y, shuffle the data, and

check the shape of each array.

```
In [25]: X_train = np.array(train_df.image_array.values.tolist())
y_train = np.asarray(train_df.target)

X_test = np.array(test_df.image_array.values.tolist())
y_test = np.asarray(test_df.target)

X_val = np.array(val_df.image_array.values.tolist())
y_val = np.asarray(val_df.target)

# helper function to shuffle data/target in the same way
def shuffle_arrays(data, target):
    assert len(data) == len(target)
    p = np.random.permutation(len(data))
    return data[p], target[p]
# shuffle all data
X_train, y_train = shuffle_arrays(X_train, y_train)
X_val, y_val = shuffle_arrays(X_val, y_val)
X_test, y_test = shuffle_arrays(X_test, y_test)

# verify shapes are accurate
print(f"Train data/target shapes: {X_train.shape}, {y_train.shape}")
print(f"Test data/target shapes: {X_test.shape}, {y_test.shape}")
print(f"Val data/target shapes: {X_val.shape}, {y_val.shape}")
```

Train data/target shapes: (5158, 256, 256), (5158,)

Test data/target shapes: (624, 256, 256), (624,)

Val data/target shapes: (16, 256, 256), (16,)

```
In [ ]: # set the data path
data_path = f"./data/"
# list of variables we are saving
datasets = ['X_train', 'y_train', 'X_test', 'y_test', 'X_val', 'y_val']

for dataset in datasets:
    np.save(file=f'{data_path}{dataset}.npy', arr=globals()[dataset])
```

Data Loading

```
In [26]: # set the data path and get the list of files
data_path = f"./data/"
data_files = list(listdir(data_path))

for file in data_files:
    # split off the components of the file name
    dataset = file[:-4]
    # set a global variable for each file with that files name and the data from the fi
    globals()[f'{dataset}'] = np.load(data_path + file)

# set the data types to float
X_train_orig = X_train.astype('float32')
X_val_orig = X_val.astype('float32')
X_test_orig = X_test.astype('float32')
```

```
# normalize the data
X_train_orig /= 255.
X_val_orig /= 255.
X_test_orig /= 255.

# verify shapes are accurate
print(f"Train data/target shapes: {X_train_orig.shape}, {y_train.shape}")
print(f"Val data/target shapes: {X_val_orig.shape}, {y_val.shape}")
print(f"Test data/target shapes: {X_test_orig.shape}, {y_test.shape}")
```

Train data/target shapes: (5158, 256, 256), (5158,)

Val data/target shapes: (16, 256, 256), (16,)

Test data/target shapes: (624, 256, 256), (624,)

Modeling helpers

Global Variables

In [27]:

```
# setting up training data class imbalance for eventual weighting
neg, pos = np.bincount(y_train)
total = neg+pos

percent_neg = round(neg/total*100, 2)
percent_pos = round(pos/total*100, 2)

print(f'Training data contains:\nTotal: {total}\nNegative: {neg} ({percent_neg}%)\nPosi
```

Training data contains:

Total: 5158

Negative: 1341 (26.0%)

Positive: 3817 (74.0%)

In [28]:

```
# setting up metrics for model comparison
METRICS = [
    keras.metrics.TruePositives(name='tp'),
    keras.metrics.FalsePositives(name='fp'),
    keras.metrics.TrueNegatives(name='tn'),
    keras.metrics.FalseNegatives(name='fn'),
    keras.metrics.BinaryAccuracy(name='accuracy'), # using binary accuracy
    keras.metrics.Precision(name='precision'),
    keras.metrics.Recall(name='recall'),
    keras.metrics.AUC(name='auc'),
    keras.metrics.AUC(name='prc', curve='PR'), # precision-recall curve
]
```

In [29]:

```
# set up the class weights using y_train
classes = list(np.unique(y_train))
weights = list(compute_class_weight(class_weight='balanced', classes=np.unique(y_train))
class_weights = dict(zip(classes, weights))
class_weights
```

Out[29]: {0: 1.9231916480238629, 1: 0.675661514278229}

Helper Functions

In [30]:

```
def set_x_y():
    # get user input for testing dataset
    print("Choose which testing dataset to use for reporting")
    print("Input S for training data\nInput T for test data\nInput V for validation dat")
    test_choice = input()
    # set the correct global datasets to the local X_test and y_test for reporting
    if test_choice.lower() == 's':
        X_test, y_test = globals()['X_train'], globals()['y_train']
    elif test_choice.lower() == 't':
        X_test, y_test = globals()['X_test'], globals()['y_test']
    else:
        X_test, y_test = globals()['X_val'], globals()['y_val']
    # return the test target and label
    return X_test, y_test

def converted_ypred(model, X_test, threshold=0.5):
    # generate predictions
    y_pred = model.predict(X_test)
    # convert the predictions based on the threshold
    # default is: 0.5
    converted_ypred = []
    for i in range(len(y_pred)):
        if y_pred[i][0] > threshold:
            converted_ypred.append(1)
        else:
            converted_ypred.append(0)
    # return the converted y_preds as an array
    return np.asarray(converted_ypred)

def model_report(y_test, y_pred):
    # set up labels for classification report and confusion matrix
    report_labels = ['normal', 'pneumonia']
    column_labels = ['predicted normal', 'predicted pneumonia']
    index_labels = ['actual normal', 'actual pneumonia']
    # generate the confusion matrix
    cmatrix = confusion_matrix(y_test, y_pred)
    # convert into dataframe
    cmatrix_df = pd.DataFrame(cmatrix, columns=column_labels, index=index_labels)
    print('-----\n')
    # print the classification report
    print(classification_report(y_test, y_pred, zero_division=0, target_names=report_labels))
    # show report as heatmap
    fig, ax = plt.subplots(figsize=(8,6))
    ax = sns.heatmap(data=cmatrix_df, annot=True, cmap='Blues', fmt='g')
    locs, labels = plt.xticks()
    plt.setp(labels, rotation=0)
    plt.show()

def plot_metrics(results):
    # set the history
    history = results.history
    # set up metrics to plot
    metrics = ['loss', 'accuracy', 'precision', 'recall']
    # set up x-axis (number of epochs)
    epochs = list(range(1, len(history['loss'])+1))
    # generate plot for each of the 4 metrics in a 2x2 plot
    plt.figure(figsize=(15,15))
```

```

for n, metric in enumerate(metrics):
    name = metric.replace("_", " ").capitalize()
    plt.subplot(2,2,n+1)
    plt.plot(epochs, history[metric], color=colors[0], label='Train')
    plt.plot(epochs, history['val_'+metric], color=colors[1], linestyle="--", label='Val')
    plt.xlabel('Epoch')
    plt.ylabel(name)
    plt.ylim([0,1.1])
plt.legend();

def full_report(model, results=None):
    # set the X_train and y_train
    X_train, y_train = globals()['X_train'], globals()['y_train']
    # set the testing X and y based on input
    X_test, y_test = set_x_y()
    # create the y_pred
    y_pred = converted_ypred(model, X_test)
    # print the model report and confusion matrix
    model_report(y_test, y_pred)
    # display the 2x2 plot of metrics
    if results != None:
        plot_metrics(results)

```

Baseline model

Correcting data shape

```

In [31]: # set the lengths of each set of data
train_len = X_train_orig.shape[0]
test_len = X_test_orig.shape[0]
val_len = X_val_orig.shape[0]

# set the new reshape size for the images
img_size = X_train_orig.shape[1] ** 2

# reshape all three data groups
X_train = X_train_orig.reshape(train_len, img_size).astype('float32')
X_test = X_test_orig.reshape(test_len, img_size).astype('float32')
X_val = X_val_orig.reshape(val_len, img_size).astype('float32')

# verify shapes are accurate
print(f"Train data/target shapes: {X_train.shape}, {y_train.shape}")
print(f"Test data/target shapes: {X_test.shape}, {y_test.shape}")
print(f"Val data/target shapes: {X_val.shape}, {y_val.shape}")

```

```

Train data/target shapes: (5158, 65536), (5158,)
Test data/target shapes: (624, 65536), (624,)
Val data/target shapes: (16, 65536), (16,)

```

MLP model

Inspiration from TensorFlow Tutorial to enact some changes to process and account for imbalanced data. [Source](#)

```

In [32]: def make_baseline_model():

```

```

# set up model
model = keras.Sequential()
model.add(Dense(128, activation='relu', input_shape = (X_train.shape[-1],)))
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

# compile model
model.compile(loss='binary_crossentropy',
              optimizer='Adam',
              metrics=METRICS)

return model

```

In [33]:

```

baseline_model = make_baseline_model()
baseline_model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 128)	8388736
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 32)	2080
dropout_2 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 1)	33

```

=====
Total params: 8,399,105
Trainable params: 8,399,105
Non-trainable params: 0

```

In [34]:

```

EPOCHS = 5
BATCH_SIZE = 30

baseline_results = baseline_model.fit(X_train, y_train,
                                     batch_size=BATCH_SIZE,
                                     epochs=EPOCHS,
                                     class_weight=dict(enumerate(class_weights)),
                                     validation_data=(X_val, y_val))

```

Epoch 1/5

```

172/172 [=====] - 8s 39ms/step - loss: 0.7830 - tp: 3650.0000 -
fp: 1300.0000 - tn: 41.0000 - fn: 167.0000 - accuracy: 0.7156 - precision: 0.7374 - reca
ll: 0.9562 - auc: 0.4954 - prc: 0.7382 - val_loss: 174.4564 - val_tp: 8.0000 - val_fp:

```

```
8.0000 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision:
0.5000 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
Epoch 2/5
172/172 [=====] - 6s 34ms/step - loss: 0.0826 - tp: 3801.0000 -
fp: 1336.0000 - tn: 5.0000 - fn: 16.0000 - accuracy: 0.7379 - precision: 0.7399 - recal
l: 0.9958 - auc: 0.4990 - prc: 0.7396 - val_loss: 191.2362 - val_tp: 8.0000 - val_fp: 8.
0000 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision:
0.5000 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
Epoch 3/5
172/172 [=====] - 6s 34ms/step - loss: 0.0116 - tp: 3810.0000 -
fp: 1339.0000 - tn: 2.0000 - fn: 7.0000 - accuracy: 0.7390 - precision: 0.7399 - recall:
0.9982 - auc: 0.4999 - prc: 0.7400 - val_loss: 211.7686 - val_tp: 8.0000 - val_fp: 8.000
0 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision: 0.50
00 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
Epoch 4/5
172/172 [=====] - 6s 34ms/step - loss: 0.0848 - tp: 3806.0000 -
fp: 1337.0000 - tn: 4.0000 - fn: 11.0000 - accuracy: 0.7387 - precision: 0.7400 - recal
l: 0.9971 - auc: 0.4999 - prc: 0.7400 - val_loss: 251.9705 - val_tp: 8.0000 - val_fp: 8.
0000 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision:
0.5000 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
Epoch 5/5
172/172 [=====] - 6s 34ms/step - loss: 0.0233 - tp: 3814.0000 -
fp: 1339.0000 - tn: 2.0000 - fn: 3.0000 - accuracy: 0.7398 - precision: 0.7402 - recall:
0.9992 - auc: 0.5000 - prc: 0.7400 - val_loss: 247.2593 - val_tp: 8.0000 - val_fp: 8.000
0 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision: 0.50
00 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
```

```
In [35]: full_report(baseline_model, baseline_results)
```

Choose which testing dataset to use for reporting






Input S for training data

Input T for test data

Input V for validation data:

```
t
-----
```

	precision	recall	f1-score	support
normal	0.00	0.00	0.00	234
pneumonia	0.62	1.00	0.77	390
accuracy			0.62	624
macro avg	0.31	0.50	0.38	624
weighted avg	0.39	0.62	0.48	624



This doesn't look great. Our model accuracy is at 74% and steady, and the validation is at 50%. Unsurprisingly, these are the percentages of our data that are labeled with pneumonia (74% for training, 50% for validation). It seems that the model is predicting the majority class for each image, definitely not good.

CNN models

Correcting data shape

In [36]:

```
# set up the X train/test/val for CNNs
X_train = X_train_orig.reshape(len(X_train_orig), 256, 256, 1)
X_val = X_val_orig.reshape(len(X_val_orig), 256, 256, 1)
X_test = X_test_orig.reshape(len(X_test_orig), 256, 256, 1)

# verify shapes are accurate
print(f"Train data/target shapes: {X_train.shape}, {y_train.shape}")
print(f"Test data/target shapes: {X_test.shape}, {y_test.shape}")
print(f"Val data/target shapes: {X_val.shape}, {y_val.shape}")
```

Train data/target shapes: (5158, 256, 256, 1), (5158,)

Test data/target shapes: (624, 256, 256, 1), (624,)

Val data/target shapes: (16, 256, 256, 1), (16,)

Callbacks

In [37]:

```
# we are going to employ a patience of 2 for early stopping
CALLBACKS = [
    EarlyStopping(patience=2),

    ModelCheckpoint(filepath='model.epoch{epoch:02d}-loss{val_loss:.2f}.hdf5',
                    monitor='val_loss',
                    verbose=1,
                    save_best_only=True,
                    mode='min'
                    )
]
```

Hyperparameter Defaults

To ensure we are only changing particular hyperparameters at a time, we will set defaults to be used in model building.

- The input shape is set correct for our images size (256, 256, 1).
- Default epochs will be 25, using early stopping this should be fine.
- Default batch size will be 128.
- Default layer activation will use 'relu'
- Default output activation will use 'sigmoid'
- Default loss will use binary crossentropy
- Default optimizer will use 'adam'.
- The list of metrics was defined in section 6.1

```
In [38]: # set the default input shape to (256, 256, 1)
INPUT_SHAPE = (256, 256, 1)

# set the default epoch length to 10
EPOCHS = 25

# set the default batch_size to 128
BATCH_SIZE = 128

# set the default layer activation to 'relu'
ACTIVATION = 'relu'

# set the default output activation to 'sigmoid'
OUTPUT = 'sigmoid'

# set the default loss to 'binary_crossentropy'
LOSS = 'binary_crossentropy'

# set the default optimizer to 'adam'
OPTIMIZER = 'adam'
```

CNN model v1

For the first CNN model, we will use 3 layers of convolution, starting with 32 filters and doubling each layer.

```
In [39]: def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=OPTIMIZER,
                  metrics=METRICS)

    return model

cnn_model_v1 = make_cnn_model()
cnn_model_v1.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense_4 (Dense)	(None, 256)	29491456
dense_5 (Dense)	(None, 1)	257
Total params: 29,584,385		
Trainable params: 29,584,385		
Non-trainable params: 0		

In [40]:

```
cnn_model_v1_results = cnn_model_v1.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         validation_data=(X_val, y_val))
```

Epoch 1/25

41/41 [=====] - ETA: 0s - loss: 0.4141 - tp: 3403.0000 - fp: 601.0000 - tn: 748.0000 - fn: 422.0000 - accuracy: 0.8023 - precision: 0.8499 - recall: 0.8897 - auc: 0.8461 - prc: 0.9200

Epoch 1: val_loss improved from inf to 0.53612, saving model to model.epoch01-loss0.54.hdf5

41/41 [=====] - 242s 6s/step - loss: 0.4141 - tp: 3403.0000 - fp: 601.0000 - tn: 748.0000 - fn: 422.0000 - accuracy: 0.8023 - precision: 0.8499 - recall: 0.8897 - auc: 0.8461 - prc: 0.9200 - val_loss: 0.5361 - val_tp: 8.0000 - val_fp: 5.0000 - val_tn: 3.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6875 - val_precision: 0.6154 - val_recall: 1.0000 - val_auc: 0.8906 - val_prc: 0.8977

Epoch 2/25

41/41 [=====] - ETA: 0s - loss: 0.1591 - tp: 3683.0000 - fp: 173.0000 - tn: 1168.0000 - fn: 134.0000 - accuracy: 0.9405 - precision: 0.9551 - recall: 0.9649 - auc: 0.9794 - prc: 0.9924

Epoch 2: val_loss did not improve from 0.53612

41/41 [=====] - 258s 6s/step - loss: 0.1591 - tp: 3683.0000 - fp: 173.0000 - tn: 1168.0000 - fn: 134.0000 - accuracy: 0.9405 - precision: 0.9551 - recall: 0.9649 - auc: 0.9794 - prc: 0.9924 - val_loss: 0.5868 - val_tp: 7.0000 - val_fp: 2.0000 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - val_recall: 0.8750 - val_auc: 0.8203 - val_prc: 0.6953

Epoch 3/25

41/41 [=====] - ETA: 0s - loss: 0.1239 - tp: 3690.0000 - fp: 124.0000 - tn: 1217.0000 - fn: 127.0000 - accuracy: 0.9513 - precision: 0.9675 - recall:

0.9667 - auc: 0.9876 - prc: 0.9953

Epoch 3: val_loss did not improve from 0.53612

41/41 [=====] - 240s 6s/step - loss: 0.1239 - tp: 3690.0000 - fp: 124.0000 - tn: 1217.0000 - fn: 127.0000 - accuracy: 0.9513 - precision: 0.9675 - recall: 0.9667 - auc: 0.9876 - prc: 0.9953 - val_loss: 0.6202 - val_tp: 8.0000 - val_fp: 5.0000 - val_tn: 3.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6875 - val_precision: 0.6154 - val_recall: 1.0000 - val_auc: 0.9531 - val_prc: 0.9493

In [41]:

```
full_report(cnn_model_v1, cnn_model_v1_results)
```

Choose which testing dataset to use for reporting


Input S for training data

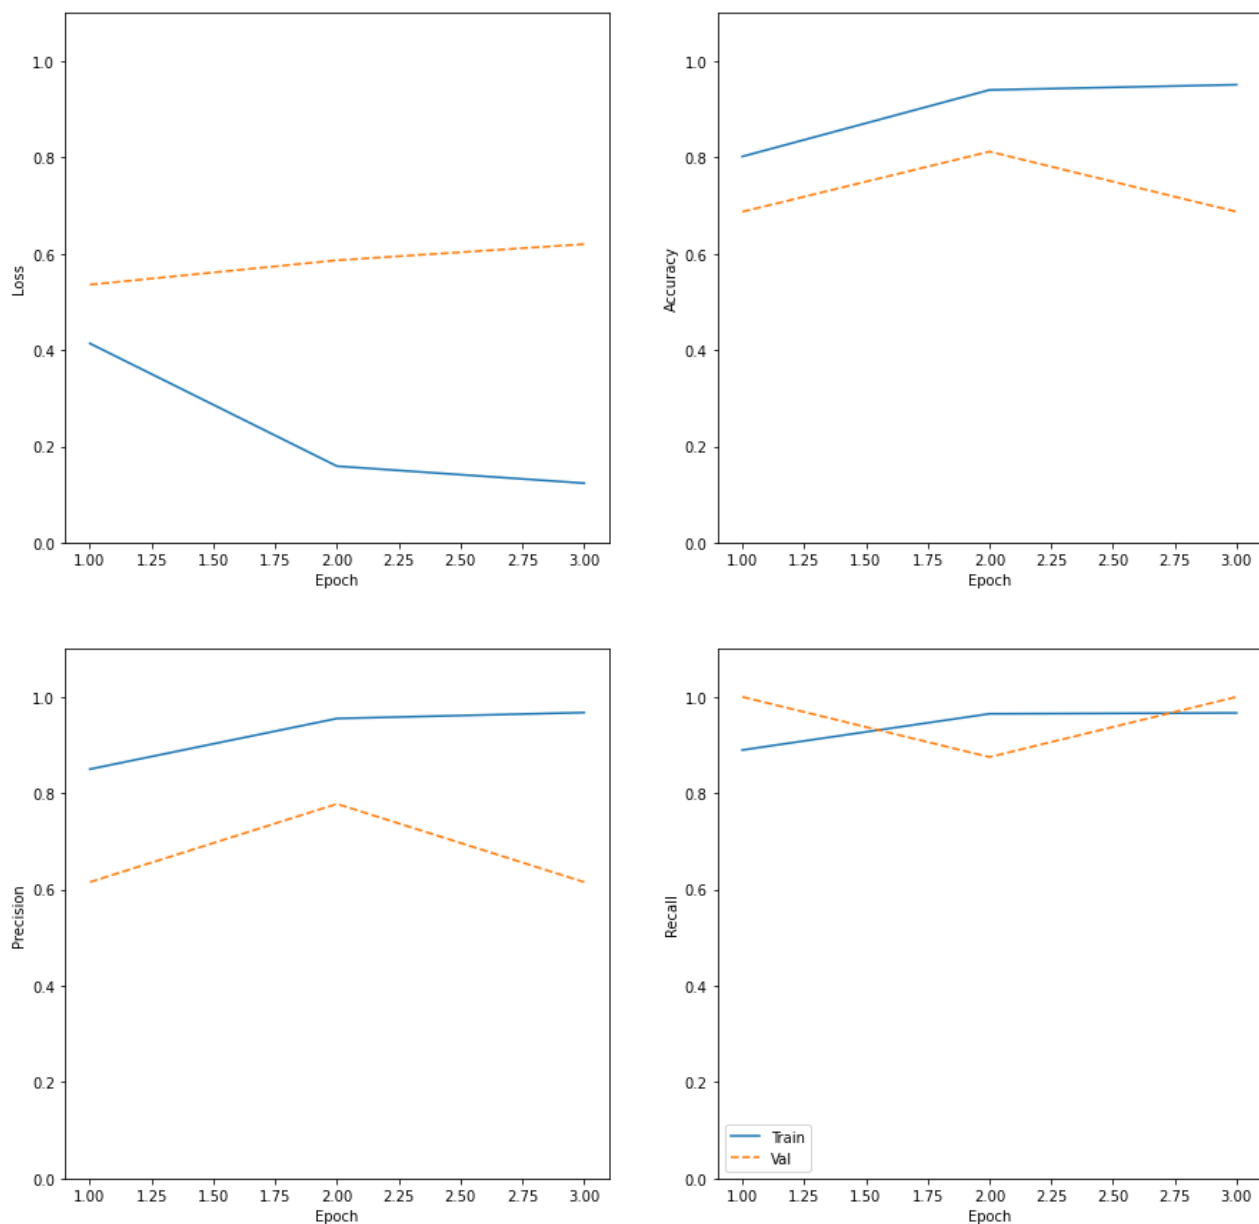
Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.97	0.27	0.43	234
pneumonia	0.70	0.99	0.82	390
accuracy			0.72	624
macro avg	0.83	0.63	0.62	624
weighted avg	0.80	0.72	0.67	624





CNN model v2

Adding another convolution layer with 64 filters

In [42]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))
```

```
model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v2 = make_cnn_model()
cnn_model_v2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_3 (MaxPooling 2D)	(None, 127, 127, 32)	0
conv2d_4 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_4 (MaxPooling 2D)	(None, 62, 62, 64)	0
conv2d_5 (Conv2D)	(None, 60, 60, 64)	36928
max_pooling2d_5 (MaxPooling 2D)	(None, 30, 30, 64)	0
conv2d_6 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_6 (MaxPooling 2D)	(None, 14, 14, 128)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_6 (Dense)	(None, 256)	6422784
dense_7 (Dense)	(None, 1)	257
=====		
Total params: 6,552,641		
Trainable params: 6,552,641		
Non-trainable params: 0		

In [43]:

```
cnn_model_v2_results = cnn_model_v2.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         validation_data=(X_val, y_val))
```

Epoch 1/25
41/41 [=====] - ETA: 0s - loss: 0.4241 - tp: 3645.0000 - fp: 78

```

6.0000 - tn: 563.0000 - fn: 180.0000 - accuracy: 0.8133 - precision: 0.8226 - recall: 0.
9529 - auc: 0.8159 - prc: 0.9242
Epoch 1: val_loss did not improve from 0.53612
41/41 [=====] - 225s 5s/step - loss: 0.4241 - tp: 3645.0000 - f
p: 786.0000 - tn: 563.0000 - fn: 180.0000 - accuracy: 0.8133 - precision: 0.8226 - recal
l: 0.9529 - auc: 0.8159 - prc: 0.9242 - val_loss: 0.6760 - val_tp: 7.0000 - val_fp: 5.00
00 - val_tn: 3.0000 - val_fn: 1.0000 - val_accuracy: 0.6250 - val_precision: 0.5833 - va
l_recall: 0.8750 - val_auc: 0.9219 - val_prc: 0.9500
Epoch 2/25
41/41 [=====] - ETA: 0s - loss: 0.1524 - tp: 3666.0000 - fp: 15
7.0000 - tn: 1184.0000 - fn: 151.0000 - accuracy: 0.9403 - precision: 0.9589 - recall:
0.9604 - auc: 0.9805 - prc: 0.9926
Epoch 2: val_loss did not improve from 0.53612
41/41 [=====] - 245s 6s/step - loss: 0.1524 - tp: 3666.0000 - f
p: 157.0000 - tn: 1184.0000 - fn: 151.0000 - accuracy: 0.9403 - precision: 0.9589 - reca
ll: 0.9604 - auc: 0.9805 - prc: 0.9926 - val_loss: 0.8186 - val_tp: 8.0000 - val_fp: 6.0
000 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714
- val_recall: 1.0000 - val_auc: 0.9141 - val_prc: 0.9101
Epoch 3/25
41/41 [=====] - ETA: 0s - loss: 0.1174 - tp: 3706.0000 - fp: 11
4.0000 - tn: 1227.0000 - fn: 111.0000 - accuracy: 0.9564 - precision: 0.9702 - recall:
0.9709 - auc: 0.9885 - prc: 0.9957
Epoch 3: val_loss did not improve from 0.53612
41/41 [=====] - 259s 6s/step - loss: 0.1174 - tp: 3706.0000 - f
p: 114.0000 - tn: 1227.0000 - fn: 111.0000 - accuracy: 0.9564 - precision: 0.9702 - reca
ll: 0.9709 - auc: 0.9885 - prc: 0.9957 - val_loss: 0.5425 - val_tp: 7.0000 - val_fp: 2.0
000 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - v
al_recall: 0.8750 - val_auc: 0.8672 - val_prc: 0.8577
Epoch 4/25
41/41 [=====] - ETA: 0s - loss: 0.1073 - tp: 3724.0000 - fp: 10
9.0000 - tn: 1232.0000 - fn: 93.0000 - accuracy: 0.9608 - precision: 0.9716 - recall: 0.
9756 - auc: 0.9903 - prc: 0.9964
Epoch 4: val_loss did not improve from 0.53612
41/41 [=====] - 238s 6s/step - loss: 0.1073 - tp: 3724.0000 - f
p: 109.0000 - tn: 1232.0000 - fn: 93.0000 - accuracy: 0.9608 - precision: 0.9716 - recal
l: 0.9756 - auc: 0.9903 - prc: 0.9964 - val_loss: 0.5925 - val_tp: 8.0000 - val_fp: 6.00
00 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714
- val_recall: 1.0000 - val_auc: 0.9219 - val_prc: 0.9283
Epoch 5/25
41/41 [=====] - ETA: 0s - loss: 0.0850 - tp: 3733.0000 - fp: 8
2.0000 - tn: 1259.0000 - fn: 84.0000 - accuracy: 0.9678 - precision: 0.9785 - recall: 0.
9780 - auc: 0.9938 - prc: 0.9977
Epoch 5: val_loss did not improve from 0.53612
41/41 [=====] - 242s 6s/step - loss: 0.0850 - tp: 3733.0000 - f
p: 82.0000 - tn: 1259.0000 - fn: 84.0000 - accuracy: 0.9678 - precision: 0.9785 - recal
l: 0.9780 - auc: 0.9938 - prc: 0.9977 - val_loss: 1.0964 - val_tp: 8.0000 - val_fp: 6.00
00 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714
- val_recall: 1.0000 - val_auc: 0.8125 - val_prc: 0.8634

```

```
In [44]: full_report(cnn_model_v2, cnn_model_v2_results)
```

Choose which testing dataset to use for reporting

Input S for training data

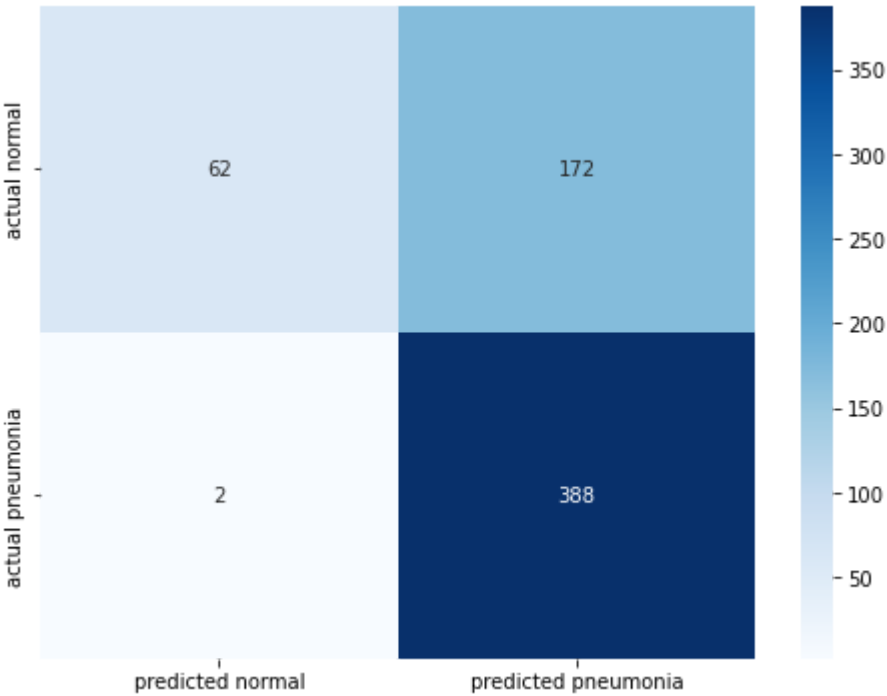
Input T for test data





Input V for validation data:

t

precision	recall	f1-score	support
-----------	--------	----------	---------

normal	0.97	0.26	0.42	234
pneumonia	0.69	0.99	0.82	390
accuracy			0.72	624
macro avg	0.83	0.63	0.62	624
weighted avg	0.80	0.72	0.67	624





CNN model v3

Adding in class weighting (calculated in 6.1)

In [45]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))
```

```
model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v3 = make_cnn_model()
cnn_model_v3.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_7 (MaxPooling 2D)	(None, 127, 127, 32)	0
conv2d_8 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_8 (MaxPooling 2D)	(None, 62, 62, 64)	0
conv2d_9 (Conv2D)	(None, 60, 60, 64)	36928
max_pooling2d_9 (MaxPooling 2D)	(None, 30, 30, 64)	0
conv2d_10 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_10 (MaxPoolin g2D)	(None, 14, 14, 128)	0
flatten_2 (Flatten)	(None, 25088)	0
dense_8 (Dense)	(None, 256)	6422784
dense_9 (Dense)	(None, 1)	257
=====		
Total params: 6,552,641		
Trainable params: 6,552,641		
Non-trainable params: 0		

In [46]:

```
cnn_model_v3_results = cnn_model_v3.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         class_weight=class_weights,
                                         validation_data=(X_val, y_val))
```

Epoch 1/25

```

41/41 [=====] - ETA: 0s - loss: 0.4961 - tp: 2844.0000 - fp: 32
9.0000 - tn: 1020.0000 - fn: 981.0000 - accuracy: 0.7468 - precision: 0.8963 - recall:
0.7435 - auc: 0.8358 - prc: 0.9314
Epoch 1: val_loss improved from 0.53612 to 0.29326, saving model to model.epoch01-loss0.
29.hdf5
41/41 [=====] - 184s 4s/step - loss: 0.4961 - tp: 2844.0000 - f
p: 329.0000 - tn: 1020.0000 - fn: 981.0000 - accuracy: 0.7468 - precision: 0.8963 - reca
ll: 0.7435 - auc: 0.8358 - prc: 0.9314 - val_loss: 0.2933 - val_tp: 8.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.9375 - val_precision: 0.8889
- val_recall: 1.0000 - val_auc: 0.9844 - val_prc: 0.9853
Epoch 2/25
41/41 [=====] - ETA: 0s - loss: 0.1611 - tp: 3564.0000 - fp: 8
9.0000 - tn: 1252.0000 - fn: 253.0000 - accuracy: 0.9337 - precision: 0.9756 - recall:
0.9337 - auc: 0.9840 - prc: 0.9939
Epoch 2: val_loss did not improve from 0.29326
41/41 [=====] - 184s 4s/step - loss: 0.1611 - tp: 3564.0000 - f
p: 89.0000 - tn: 1252.0000 - fn: 253.0000 - accuracy: 0.9337 - precision: 0.9756 - recal
l: 0.9337 - auc: 0.9840 - prc: 0.9939 - val_loss: 0.4653 - val_tp: 8.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.9219 - val_prc: 0.9283
Epoch 3/25
41/41 [=====] - ETA: 0s - loss: 0.1292 - tp: 3635.0000 - fp: 6
1.0000 - tn: 1280.0000 - fn: 182.0000 - accuracy: 0.9529 - precision: 0.9835 - recall:
0.9523 - auc: 0.9889 - prc: 0.9959
Epoch 3: val_loss did not improve from 0.29326
41/41 [=====] - 182s 4s/step - loss: 0.1292 - tp: 3635.0000 - f
p: 61.0000 - tn: 1280.0000 - fn: 182.0000 - accuracy: 0.9529 - precision: 0.9835 - recal
l: 0.9523 - auc: 0.9889 - prc: 0.9959 - val_loss: 0.3937 - val_tp: 6.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 2.0000 - val_accuracy: 0.7500 - val_precision: 0.7500 - va
l_recall: 0.7500 - val_auc: 0.9062 - val_prc: 0.9108

```

In [47]:

```
full_report(cnn_model_v3, cnn_model_v3_results)
```

Choose which testing dataset to use for reporting





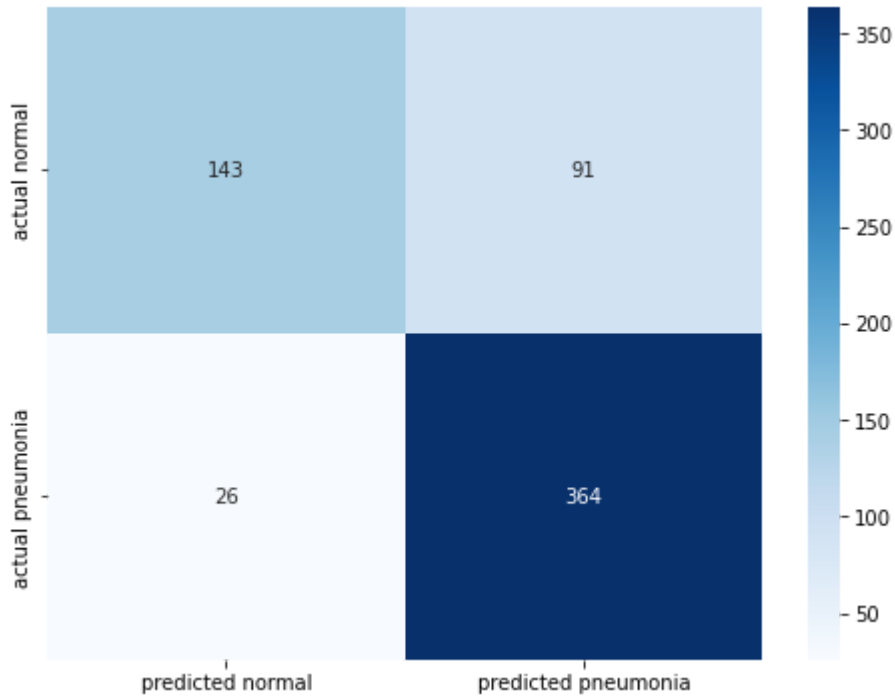
Input S for training data

Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.85	0.61	0.71	234
pneumonia	0.80	0.93	0.86	390
accuracy			0.81	624
macro avg	0.82	0.77	0.79	624
weighted avg	0.82	0.81	0.80	624



CNN Model v4

Change the optimizer to RMSProp with a learning rate of 0.01

In [48]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=keras.optimizers.RMSprop(0.01),
                  metrics=METRICS)

    return model

cnn_model_v4 = make_cnn_model()
cnn_model_v4.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_11 (MaxPoolin g2D)	(None, 127, 127, 32)	0
conv2d_12 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_12 (MaxPoolin g2D)	(None, 62, 62, 64)	0
conv2d_13 (Conv2D)	(None, 60, 60, 64)	36928
max_pooling2d_13 (MaxPoolin g2D)	(None, 30, 30, 64)	0
conv2d_14 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_14 (MaxPoolin	(None, 14, 14, 128)	0

g2D)

flatten_3 (Flatten)	(None, 25088)	0
dense_10 (Dense)	(None, 256)	6422784
dense_11 (Dense)	(None, 1)	257

```

=====
Total params: 6,552,641
Trainable params: 6,552,641
Non-trainable params: 0

```

In [49]:

```

cnn_model_v4_results = cnn_model_v4.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         class_weight=class_weights,
                                         validation_data=(X_val, y_val))

```

Epoch 1/25

```

41/41 [=====] - ETA: 0s - loss: 583.2070 - tp: 2322.0000 - fp:
436.0000 - tn: 913.0000 - fn: 1503.0000 - accuracy: 0.6252 - precision: 0.8419 - recall:
0.6071 - auc: 0.6558 - prc: 0.8140

```

Epoch 1: val_loss did not improve from 0.29326

```

41/41 [=====] - 181s 4s/step - loss: 583.2070 - tp: 2322.0000 -
fp: 436.0000 - tn: 913.0000 - fn: 1503.0000 - accuracy: 0.6252 - precision: 0.8419 - rec
all: 0.6071 - auc: 0.6558 - prc: 0.8140 - val_loss: 1.0111 - val_tp: 0.0000e+00 - val_f
p: 0.0000e+00 - val_tn: 8.0000 - val_fn: 8.0000 - val_accuracy: 0.5000 - val_precision:
0.0000e+00 - val_recall: 0.0000e+00 - val_auc: 0.8125 - val_prc: 0.8461

```

Epoch 2/25

```

41/41 [=====] - ETA: 0s - loss: 1.6290 - tp: 2732.0000 - fp: 38
8.0000 - tn: 953.0000 - fn: 1085.0000 - accuracy: 0.7144 - precision: 0.8756 - recall:
0.7157 - auc: 0.7574 - prc: 0.8623

```

Epoch 2: val_loss did not improve from 0.29326

```

41/41 [=====] - 187s 5s/step - loss: 1.6290 - tp: 2732.0000 - f
p: 388.0000 - tn: 953.0000 - fn: 1085.0000 - accuracy: 0.7144 - precision: 0.8756 - reca
ll: 0.7157 - auc: 0.7574 - prc: 0.8623 - val_loss: 0.6657 - val_tp: 8.0000 - val_fp: 7.0
000 - val_tn: 1.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.5625 - val_precision: 0.5333
- val_recall: 1.0000 - val_auc: 0.8750 - val_prc: 0.9085

```

Epoch 3/25

```

41/41 [=====] - ETA: 0s - loss: 2.0830 - tp: 2544.0000 - fp: 50
7.0000 - tn: 834.0000 - fn: 1273.0000 - accuracy: 0.6549 - precision: 0.8338 - recall:
0.6665 - auc: 0.6456 - prc: 0.7960

```

Epoch 3: val_loss did not improve from 0.29326

```

41/41 [=====] - 180s 4s/step - loss: 2.0830 - tp: 2544.0000 - f
p: 507.0000 - tn: 834.0000 - fn: 1273.0000 - accuracy: 0.6549 - precision: 0.8338 - reca
ll: 0.6665 - auc: 0.6456 - prc: 0.7960 - val_loss: 0.4074 - val_tp: 8.0000 - val_fp: 2.0
000 - val_tn: 6.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8750 - val_precision: 0.8000
- val_recall: 1.0000 - val_auc: 0.9844 - val_prc: 0.9853

```

Epoch 4/25

```

41/41 [=====] - ETA: 0s - loss: 11.2761 - tp: 2889.0000 - fp: 5
84.0000 - tn: 757.0000 - fn: 928.0000 - accuracy: 0.7069 - precision: 0.8318 - recall:
0.7569 - auc: 0.6840 - prc: 0.8182

```

Epoch 4: val_loss did not improve from 0.29326

```

41/41 [=====] - 180s 4s/step - loss: 11.2761 - tp: 2889.0000 -
fp: 584.0000 - tn: 757.0000 - fn: 928.0000 - accuracy: 0.7069 - precision: 0.8318 - reca
ll: 0.7569 - auc: 0.6840 - prc: 0.8182 - val_loss: 0.5929 - val_tp: 8.0000 - val_fp: 7.0

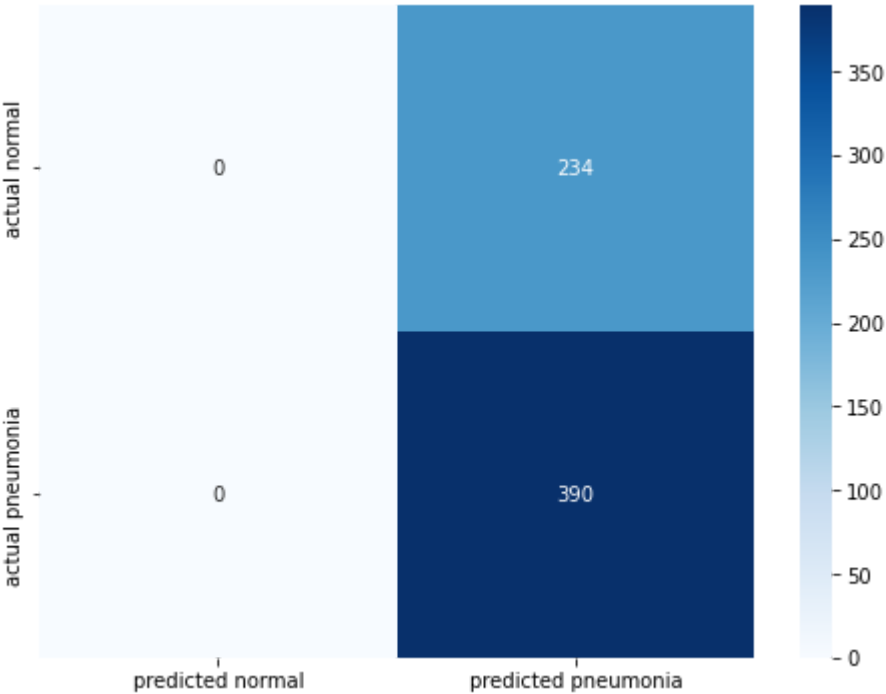
```





```
000 - val_tn: 1.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.5625 - val_precision: 0.5333
- val_recall: 1.0000 - val_auc: 1.0000 - val_prc: 1.0000
Epoch 5/25
41/41 [=====] - ETA: 0s - loss: 0.8408 - tp: 3176.0000 - fp: 48
5.0000 - tn: 856.0000 - fn: 641.0000 - accuracy: 0.7817 - precision: 0.8675 - recall: 0.
8321 - auc: 0.7799 - prc: 0.8713
Epoch 5: val_loss did not improve from 0.29326
41/41 [=====] - 174s 4s/step - loss: 0.8408 - tp: 3176.0000 - f
p: 485.0000 - tn: 856.0000 - fn: 641.0000 - accuracy: 0.7817 - precision: 0.8675 - recal
l: 0.8321 - auc: 0.7799 - prc: 0.8713 - val_loss: 10.5880 - val_tp: 8.0000 - val_fp: 8.0
000 - val_tn: 0.0000e+00 - val_fn: 0.0000e+00 - val_accuracy: 0.5000 - val_precision: 0.
5000 - val_recall: 1.0000 - val_auc: 0.5000 - val_prc: 0.5000
```

```
In [50]: full_report(cnn_model_v4, cnn_model_v4_results)
```

Choose which testing dataset to use for reporting
Input S for training data
Input T for test data
Input V for validation data:
t

	precision	recall	f1-score	support
normal	0.00	0.00	0.00	234
pneumonia	0.62	1.00	0.77	390
accuracy			0.62	624
macro avg	0.31	0.50	0.38	624
weighted avg	0.39	0.62	0.48	624





CNN model v5

Decrease the learning rate by a factor of 10.

In [51]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))
```

```
model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=keras.optimizers.RMSprop(0.001),
              metrics=METRICS)

return model

cnn_model_v5 = make_cnn_model()
cnn_model_v5.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_15 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_15 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_16 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_16 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_17 (Conv2D)	(None, 60, 60, 64)	36928
max_pooling2d_17 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_18 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_18 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten_4 (Flatten)	(None, 25088)	0
dense_12 (Dense)	(None, 256)	6422784
dense_13 (Dense)	(None, 1)	257
=====		
Total params: 6,552,641		
Trainable params: 6,552,641		
Non-trainable params: 0		

In [52]:

```
cnn_model_v5_results = cnn_model_v5.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         class_weight=class_weights,
                                         validation_data=(X_val, y_val))
```

Epoch 1/25

41/41 [=====] - ETA: 0s - loss: 0.7913 - tp: 2250.0000 - fp: 409.0000 - tn: 940.0000 - fn: 1575.0000 - accuracy: 0.6165 - precision: 0.8462 - recall: 0.5882 - auc: 0.6969 - prc: 0.8437
Epoch 1: val_loss did not improve from 0.29326
41/41 [=====] - 183s 4s/step - loss: 0.7913 - tp: 2250.0000 - fp: 409.0000 - tn: 940.0000 - fn: 1575.0000 - accuracy: 0.6165 - precision: 0.8462 - recall: 0.5882 - auc: 0.6969 - prc: 0.8437 - val_loss: 0.4092 - val_tp: 8.0000 - val_fp: 4.0000 - val_tn: 4.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.7500 - val_precision: 0.6667 - val_recall: 1.0000 - val_auc: 0.9531 - val_prc: 0.9643
Epoch 2/25
41/41 [=====] - ETA: 0s - loss: 0.3409 - tp: 3316.0000 - fp: 185.0000 - tn: 1156.0000 - fn: 501.0000 - accuracy: 0.8670 - precision: 0.9472 - recall: 0.8687 - auc: 0.9323 - prc: 0.9678
Epoch 2: val_loss did not improve from 0.29326
41/41 [=====] - 180s 4s/step - loss: 0.3409 - tp: 3316.0000 - fp: 185.0000 - tn: 1156.0000 - fn: 501.0000 - accuracy: 0.8670 - precision: 0.9472 - recall: 0.8687 - auc: 0.9323 - prc: 0.9678 - val_loss: 0.3299 - val_tp: 7.0000 - val_fp: 1.0000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - val_recall: 0.8750 - val_auc: 0.9375 - val_prc: 0.9565
Epoch 3/25
41/41 [=====] - ETA: 0s - loss: 0.2178 - tp: 3474.0000 - fp: 108.0000 - tn: 1233.0000 - fn: 343.0000 - accuracy: 0.9126 - precision: 0.9698 - recall: 0.9101 - auc: 0.9711 - prc: 0.9887
Epoch 3: val_loss did not improve from 0.29326
41/41 [=====] - 181s 4s/step - loss: 0.2178 - tp: 3474.0000 - fp: 108.0000 - tn: 1233.0000 - fn: 343.0000 - accuracy: 0.9126 - precision: 0.9698 - recall: 0.9101 - auc: 0.9711 - prc: 0.9887 - val_loss: 0.8164 - val_tp: 8.0000 - val_fp: 6.0000 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714 - val_recall: 1.0000 - val_auc: 0.9688 - val_prc: 0.9737
Epoch 4/25
41/41 [=====] - ETA: 0s - loss: 0.2927 - tp: 3467.0000 - fp: 118.0000 - tn: 1223.0000 - fn: 350.0000 - accuracy: 0.9093 - precision: 0.9671 - recall: 0.9083 - auc: 0.9556 - prc: 0.9746
Epoch 4: val_loss improved from 0.29326 to 0.25460, saving model to model.epoch04-loss0.25.hdf5
41/41 [=====] - 181s 4s/step - loss: 0.2927 - tp: 3467.0000 - fp: 118.0000 - tn: 1223.0000 - fn: 350.0000 - accuracy: 0.9093 - precision: 0.9671 - recall: 0.9083 - auc: 0.9556 - prc: 0.9746 - val_loss: 0.2546 - val_tp: 6.0000 - val_fp: 0.0000e+00 - val_tn: 8.0000 - val_fn: 2.0000 - val_accuracy: 0.8750 - val_precision: 1.0000 - val_recall: 0.7500 - val_auc: 0.9844 - val_prc: 0.9853
Epoch 5/25
41/41 [=====] - ETA: 0s - loss: 0.1743 - tp: 3548.0000 - fp: 82.0000 - tn: 1259.0000 - fn: 269.0000 - accuracy: 0.9320 - precision: 0.9774 - recall: 0.9295 - auc: 0.9807 - prc: 0.9928
Epoch 5: val_loss did not improve from 0.25460
41/41 [=====] - 181s 4s/step - loss: 0.1743 - tp: 3548.0000 - fp: 82.0000 - tn: 1259.0000 - fn: 269.0000 - accuracy: 0.9320 - precision: 0.9774 - recall: 0.9295 - auc: 0.9807 - prc: 0.9928 - val_loss: 0.3517 - val_tp: 8.0000 - val_fp: 2.0000 - val_tn: 6.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8750 - val_precision: 0.8000 - val_recall: 1.0000 - val_auc: 0.9531 - val_prc: 0.9493
Epoch 6/25
41/41 [=====] - ETA: 0s - loss: 0.1363 - tp: 3595.0000 - fp: 59.0000 - tn: 1282.0000 - fn: 222.0000 - accuracy: 0.9455 - precision: 0.9839 - recall: 0.9418 - auc: 0.9878 - prc: 0.9956
Epoch 6: val_loss did not improve from 0.25460
41/41 [=====] - 178s 4s/step - loss: 0.1363 - tp: 3595.0000 - fp: 59.0000 - tn: 1282.0000 - fn: 222.0000 - accuracy: 0.9455 - precision: 0.9839 - recall: 0.9418 - auc: 0.9878 - prc: 0.9956 - val_loss: 0.3066 - val_tp: 7.0000 - val_fp: 1.0000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - val_recall: 0.8750 - val_auc: 0.9531 - val_prc: 0.9493

```
In [53]: full_report(cnn_model_v5, cnn_model_v5_results)
```

Choose which testing dataset to use for reporting

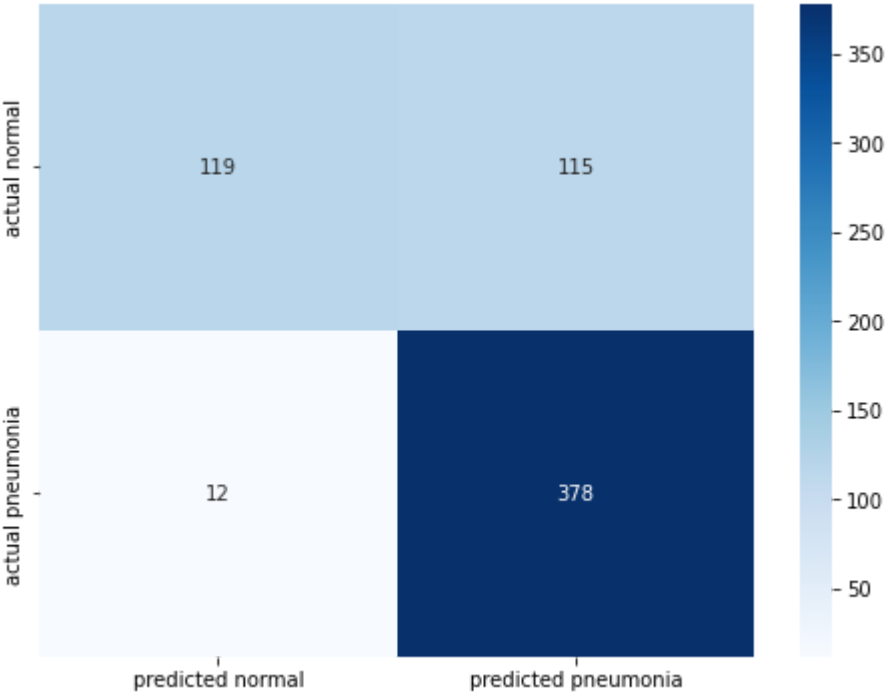
Input S for training data

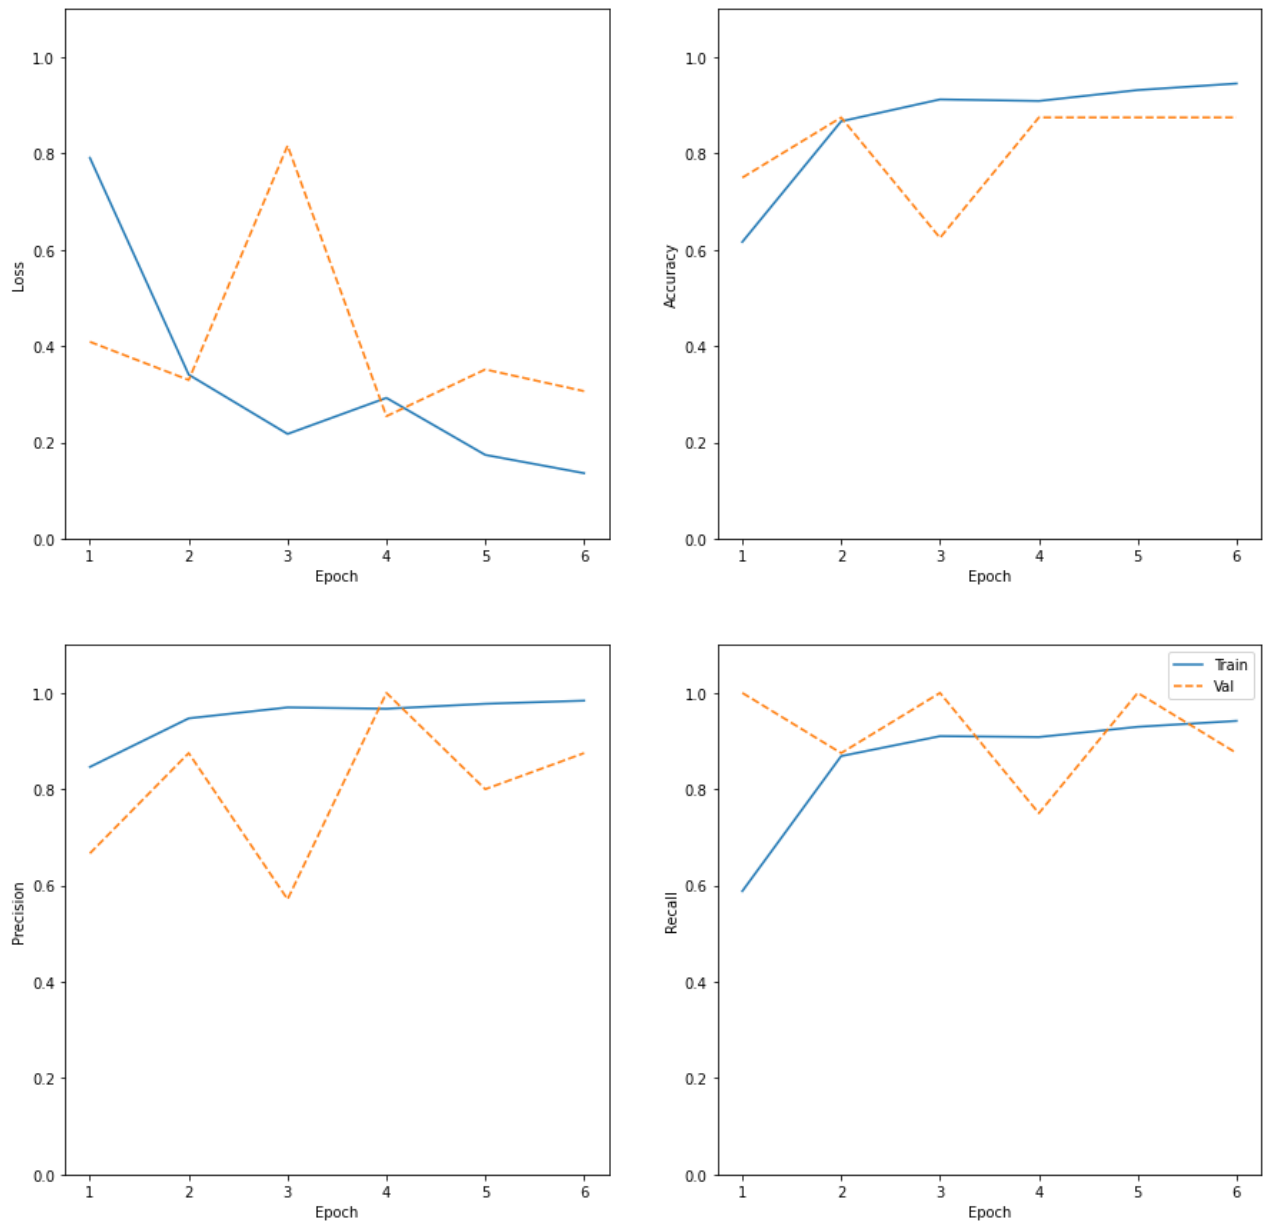
Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.91	0.51	0.65	234
pneumonia	0.77	0.97	0.86	390
accuracy			0.80	624
macro avg	0.84	0.74	0.75	624
weighted avg	0.82	0.80	0.78	624





CNN model v6

So far our implemented changes have resulted in worse performing models than our first model.

We are going to go back to that and change filters on the convolutional layers, decreasing each by half.

```
In [54]: def make_cnn_model():
    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))
```

```
model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

return model

cnn_model_v6 = make_cnn_model()
cnn_model_v6.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
=====		
conv2d_19 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_19 (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_20 (Conv2D)	(None, 125, 125, 32)	4640
max_pooling2d_20 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_21 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d_21 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten_5 (Flatten)	(None, 57600)	0
dense_14 (Dense)	(None, 256)	14745856
dense_15 (Dense)	(None, 1)	257
=====		
Total params: 14,769,409		
Trainable params: 14,769,409		
Non-trainable params: 0		

In [55]:

```
cnn_model_v6_results = cnn_model_v6.fit(X_train, y_train,
                                       batch_size=BATCH_SIZE,
                                       epochs=EPOCHS,
                                       callbacks=CALLBACKS,
                                       class_weight=class_weights,
                                       validation_data=(X_val, y_val))
```

Epoch 1/25
41/41 [=====] - ETA: 0s - loss: 0.4282 - tp: 2981.0000 - fp: 20
6.0000 - tn: 1143.0000 - fn: 844.0000 - accuracy: 0.7971 - precision: 0.9354 - recall:

0.7793 - auc: 0.8940 - prc: 0.9536
Epoch 1: val_loss did not improve from 0.25460
41/41 [=====] - 95s 2s/step - loss: 0.4282 - tp: 2981.0000 - f
p: 206.0000 - tn: 1143.0000 - fn: 844.0000 - accuracy: 0.7971 - precision: 0.9354 - reca
ll: 0.7793 - auc: 0.8940 - prc: 0.9536 - val_loss: 0.4718 - val_tp: 8.0000 - val_fp: 3.0
000 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.9219 - val_prc: 0.9283
Epoch 2/25
41/41 [=====] - ETA: 0s - loss: 0.1791 - tp: 3569.0000 - fp: 9
9.0000 - tn: 1242.0000 - fn: 248.0000 - accuracy: 0.9327 - precision: 0.9730 - recall:
0.9350 - auc: 0.9802 - prc: 0.9928
Epoch 2: val_loss did not improve from 0.25460
41/41 [=====] - 93s 2s/step - loss: 0.1791 - tp: 3569.0000 - f
p: 99.0000 - tn: 1242.0000 - fn: 248.0000 - accuracy: 0.9327 - precision: 0.9730 - recal
l: 0.9350 - auc: 0.9802 - prc: 0.9928 - val_loss: 0.3126 - val_tp: 6.0000 - val_fp: 1.00
00 - val_tn: 7.0000 - val_fn: 2.0000 - val_accuracy: 0.8125 - val_precision: 0.8571 - va
l_recall: 0.7500 - val_auc: 0.9219 - val_prc: 0.9149
Epoch 3/25
41/41 [=====] - ETA: 0s - loss: 0.1138 - tp: 3652.0000 - fp: 5
2.0000 - tn: 1289.0000 - fn: 165.0000 - accuracy: 0.9579 - precision: 0.9860 - recall:
0.9568 - auc: 0.9911 - prc: 0.9968
Epoch 3: val_loss did not improve from 0.25460
41/41 [=====] - 93s 2s/step - loss: 0.1138 - tp: 3652.0000 - f
p: 52.0000 - tn: 1289.0000 - fn: 165.0000 - accuracy: 0.9579 - precision: 0.9860 - recal
l: 0.9568 - auc: 0.9911 - prc: 0.9968 - val_loss: 0.3612 - val_tp: 8.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8750 - val_precision: 0.8000
- val_recall: 1.0000 - val_auc: 0.9375 - val_prc: 0.9265
Epoch 4/25
41/41 [=====] - ETA: 0s - loss: 0.1053 - tp: 3672.0000 - fp: 5
7.0000 - tn: 1284.0000 - fn: 145.0000 - accuracy: 0.9608 - precision: 0.9847 - recall:
0.9620 - auc: 0.9926 - prc: 0.9973
Epoch 4: val_loss did not improve from 0.25460
41/41 [=====] - 102s 2s/step - loss: 0.1053 - tp: 3672.0000 - f
p: 57.0000 - tn: 1284.0000 - fn: 145.0000 - accuracy: 0.9608 - precision: 0.9847 - recal
l: 0.9620 - auc: 0.9926 - prc: 0.9973 - val_loss: 0.2903 - val_tp: 7.0000 - val_fp: 1.00
00 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - va
l_recall: 0.8750 - val_auc: 0.9531 - val_prc: 0.9493
Epoch 5/25
41/41 [=====] - ETA: 0s - loss: 0.0917 - tp: 3688.0000 - fp: 3
6.0000 - tn: 1305.0000 - fn: 129.0000 - accuracy: 0.9680 - precision: 0.9903 - recall:
0.9662 - auc: 0.9944 - prc: 0.9979
Epoch 5: val_loss did not improve from 0.25460
41/41 [=====] - 99s 2s/step - loss: 0.0917 - tp: 3688.0000 - f
p: 36.0000 - tn: 1305.0000 - fn: 129.0000 - accuracy: 0.9680 - precision: 0.9903 - recal
l: 0.9662 - auc: 0.9944 - prc: 0.9979 - val_loss: 0.3546 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.8750 - val_prc: 0.8723
Epoch 6/25
41/41 [=====] - ETA: 0s - loss: 0.0669 - tp: 3721.0000 - fp: 3
1.0000 - tn: 1310.0000 - fn: 96.0000 - accuracy: 0.9754 - precision: 0.9917 - recall: 0.
9748 - auc: 0.9967 - prc: 0.9987
Epoch 6: val_loss did not improve from 0.25460
41/41 [=====] - 94s 2s/step - loss: 0.0669 - tp: 3721.0000 - f
p: 31.0000 - tn: 1310.0000 - fn: 96.0000 - accuracy: 0.9754 - precision: 0.9917 - recal
l: 0.9748 - auc: 0.9967 - prc: 0.9987 - val_loss: 0.2822 - val_tp: 8.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8750 - val_precision: 0.8000
- val_recall: 1.0000 - val_auc: 0.9531 - val_prc: 0.9570
Epoch 7/25
41/41 [=====] - ETA: 0s - loss: 0.0685 - tp: 3718.0000 - fp: 2
9.0000 - tn: 1312.0000 - fn: 99.0000 - accuracy: 0.9752 - precision: 0.9923 - recall: 0.

```

9741 - auc: 0.9966 - prc: 0.9989
Epoch 7: val_loss did not improve from 0.25460
41/41 [=====] - 93s 2s/step - loss: 0.0685 - tp: 3718.0000 - f
p: 29.0000 - tn: 1312.0000 - fn: 99.0000 - accuracy: 0.9752 - precision: 0.9923 - recal
l: 0.9741 - auc: 0.9966 - prc: 0.9989 - val_loss: 0.2750 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.9375 - val_prc: 0.9442
Epoch 8/25
41/41 [=====] - ETA: 0s - loss: 0.0606 - tp: 3728.0000 - fp: 2
1.0000 - tn: 1320.0000 - fn: 89.0000 - accuracy: 0.9787 - precision: 0.9944 - recall: 0.
9767 - auc: 0.9975 - prc: 0.9992
Epoch 8: val_loss did not improve from 0.25460
41/41 [=====] - 103s 3s/step - loss: 0.0606 - tp: 3728.0000 - f
p: 21.0000 - tn: 1320.0000 - fn: 89.0000 - accuracy: 0.9787 - precision: 0.9944 - recal
l: 0.9767 - auc: 0.9975 - prc: 0.9992 - val_loss: 0.4908 - val_tp: 8.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.8906 - val_prc: 0.9055
Epoch 9/25
41/41 [=====] - ETA: 0s - loss: 0.0425 - tp: 3757.0000 - fp: 2
0.0000 - tn: 1321.0000 - fn: 60.0000 - accuracy: 0.9845 - precision: 0.9947 - recall: 0.
9843 - auc: 0.9988 - prc: 0.9996
Epoch 9: val_loss did not improve from 0.25460
41/41 [=====] - 99s 2s/step - loss: 0.0425 - tp: 3757.0000 - f
p: 20.0000 - tn: 1321.0000 - fn: 60.0000 - accuracy: 0.9845 - precision: 0.9947 - recal
l: 0.9843 - auc: 0.9988 - prc: 0.9996 - val_loss: 0.5180 - val_tp: 8.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8750 - val_precision: 0.8000
- val_recall: 1.0000 - val_auc: 0.8750 - val_prc: 0.8383

```

```
In [56]: full_report(cnn_model_v6, cnn_model_v6_results)
```

Choose which testing dataset to use for reporting

Input S for training data

Input T for test data

Input V for validation data:






t

```

-----

```

	precision	recall	f1-score	support
normal	0.94	0.51	0.66	234
pneumonia	0.77	0.98	0.86	390
accuracy			0.81	624
macro avg	0.86	0.75	0.76	624
weighted avg	0.84	0.81	0.79	624



CNN model v7

Lets add in a 4th convolutional layer, duplicating the middle layer with 32 filters.

```
In [57]: def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(16, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=OPTIMIZER,
                  metrics=METRICS)

    return model

cnn_model_v7 = make_cnn_model()
cnn_model_v7.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
conv2d_22 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_22 (MaxPoolin g2D)	(None, 127, 127, 16)	0
conv2d_23 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_23 (MaxPoolin g2D)	(None, 62, 62, 16)	0
conv2d_24 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_24 (MaxPoolin g2D)	(None, 30, 30, 32)	0
conv2d_25 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_25 (MaxPoolin	(None, 14, 14, 64)	0

g2D)

flatten_6 (Flatten)	(None, 12544)	0
dense_16 (Dense)	(None, 256)	3211520
dense_17 (Dense)	(None, 1)	257

```

=====
Total params: 3,237,393
Trainable params: 3,237,393
Non-trainable params: 0

```

In [58]:

```

cnn_model_v7_results = cnn_model_v7.fit(X_train, y_train,
                                         batch_size=BATCH_SIZE,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         class_weight=class_weights,
                                         validation_data=(X_val, y_val))

```

Epoch 1/25

```

41/41 [=====] - ETA: 0s - loss: 0.4235 - tp: 2992.0000 - fp: 29
4.0000 - tn: 1055.0000 - fn: 833.0000 - accuracy: 0.7822 - precision: 0.9105 - recall:
0.7822 - auc: 0.8830 - prc: 0.9527

```

Epoch 1: val_loss did not improve from 0.25460

```

41/41 [=====] - 82s 2s/step - loss: 0.4235 - tp: 2992.0000 - f
p: 294.0000 - tn: 1055.0000 - fn: 833.0000 - accuracy: 0.7822 - precision: 0.9105 - reca
ll: 0.7822 - auc: 0.8830 - prc: 0.9527 - val_loss: 0.5510 - val_tp: 7.0000 - val_fp: 3.0
000 - val_tn: 5.0000 - val_fn: 1.0000 - val_accuracy: 0.7500 - val_precision: 0.7000 - v
al_recall: 0.8750 - val_auc: 0.8906 - val_prc: 0.9055

```

Epoch 2/25

```

41/41 [=====] - ETA: 0s - loss: 0.1928 - tp: 3513.0000 - fp: 9
1.0000 - tn: 1250.0000 - fn: 304.0000 - accuracy: 0.9234 - precision: 0.9748 - recall:
0.9204 - auc: 0.9772 - prc: 0.9914

```

Epoch 2: val_loss did not improve from 0.25460

```

41/41 [=====] - 82s 2s/step - loss: 0.1928 - tp: 3513.0000 - f
p: 91.0000 - tn: 1250.0000 - fn: 304.0000 - accuracy: 0.9234 - precision: 0.9748 - recal
l: 0.9204 - auc: 0.9772 - prc: 0.9914 - val_loss: 0.4212 - val_tp: 8.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.9219 - val_prc: 0.9249

```

Epoch 3/25

```

41/41 [=====] - ETA: 0s - loss: 0.1592 - tp: 3577.0000 - fp: 8
1.0000 - tn: 1260.0000 - fn: 240.0000 - accuracy: 0.9378 - precision: 0.9779 - recall:
0.9371 - auc: 0.9841 - prc: 0.9940

```

Epoch 3: val_loss did not improve from 0.25460

```

41/41 [=====] - 84s 2s/step - loss: 0.1592 - tp: 3577.0000 - f
p: 81.0000 - tn: 1260.0000 - fn: 240.0000 - accuracy: 0.9378 - precision: 0.9779 - recal
l: 0.9371 - auc: 0.9841 - prc: 0.9940 - val_loss: 0.3904 - val_tp: 6.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 2.0000 - val_accuracy: 0.7500 - val_precision: 0.7500 - va
l_recall: 0.7500 - val_auc: 0.8906 - val_prc: 0.9014

```

Epoch 4/25

```

41/41 [=====] - ETA: 0s - loss: 0.1192 - tp: 3641.0000 - fp: 5
8.0000 - tn: 1283.0000 - fn: 176.0000 - accuracy: 0.9546 - precision: 0.9843 - recall:
0.9539 - auc: 0.9907 - prc: 0.9965

```

Epoch 4: val_loss did not improve from 0.25460

```

41/41 [=====] - 80s 2s/step - loss: 0.1192 - tp: 3641.0000 - f
p: 58.0000 - tn: 1283.0000 - fn: 176.0000 - accuracy: 0.9546 - precision: 0.9843 - recal
l: 0.9539 - auc: 0.9907 - prc: 0.9965 - val_loss: 0.4330 - val_tp: 6.0000 - val_fp: 2.00

```

```

00 - val_tn: 6.0000 - val_fn: 2.0000 - val_accuracy: 0.7500 - val_precision: 0.7500 - va
l_recall: 0.7500 - val_auc: 0.8906 - val_prc: 0.9014
Epoch 5/25
41/41 [=====] - ETA: 0s - loss: 0.1040 - tp: 3673.0000 - fp: 5
1.0000 - tn: 1290.0000 - fn: 144.0000 - accuracy: 0.9622 - precision: 0.9863 - recall:
0.9623 - auc: 0.9926 - prc: 0.9972
Epoch 5: val_loss did not improve from 0.25460
41/41 [=====] - 81s 2s/step - loss: 0.1040 - tp: 3673.0000 - f
p: 51.0000 - tn: 1290.0000 - fn: 144.0000 - accuracy: 0.9622 - precision: 0.9863 - recal
l: 0.9623 - auc: 0.9926 - prc: 0.9972 - val_loss: 0.5486 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.8281 - val_prc: 0.8173

```

In [59]:

```
full_report(cnn_model_v7, cnn_model_v7_results)
```

Choose which testing dataset to use for reporting

Input S for training data


Input T for test data

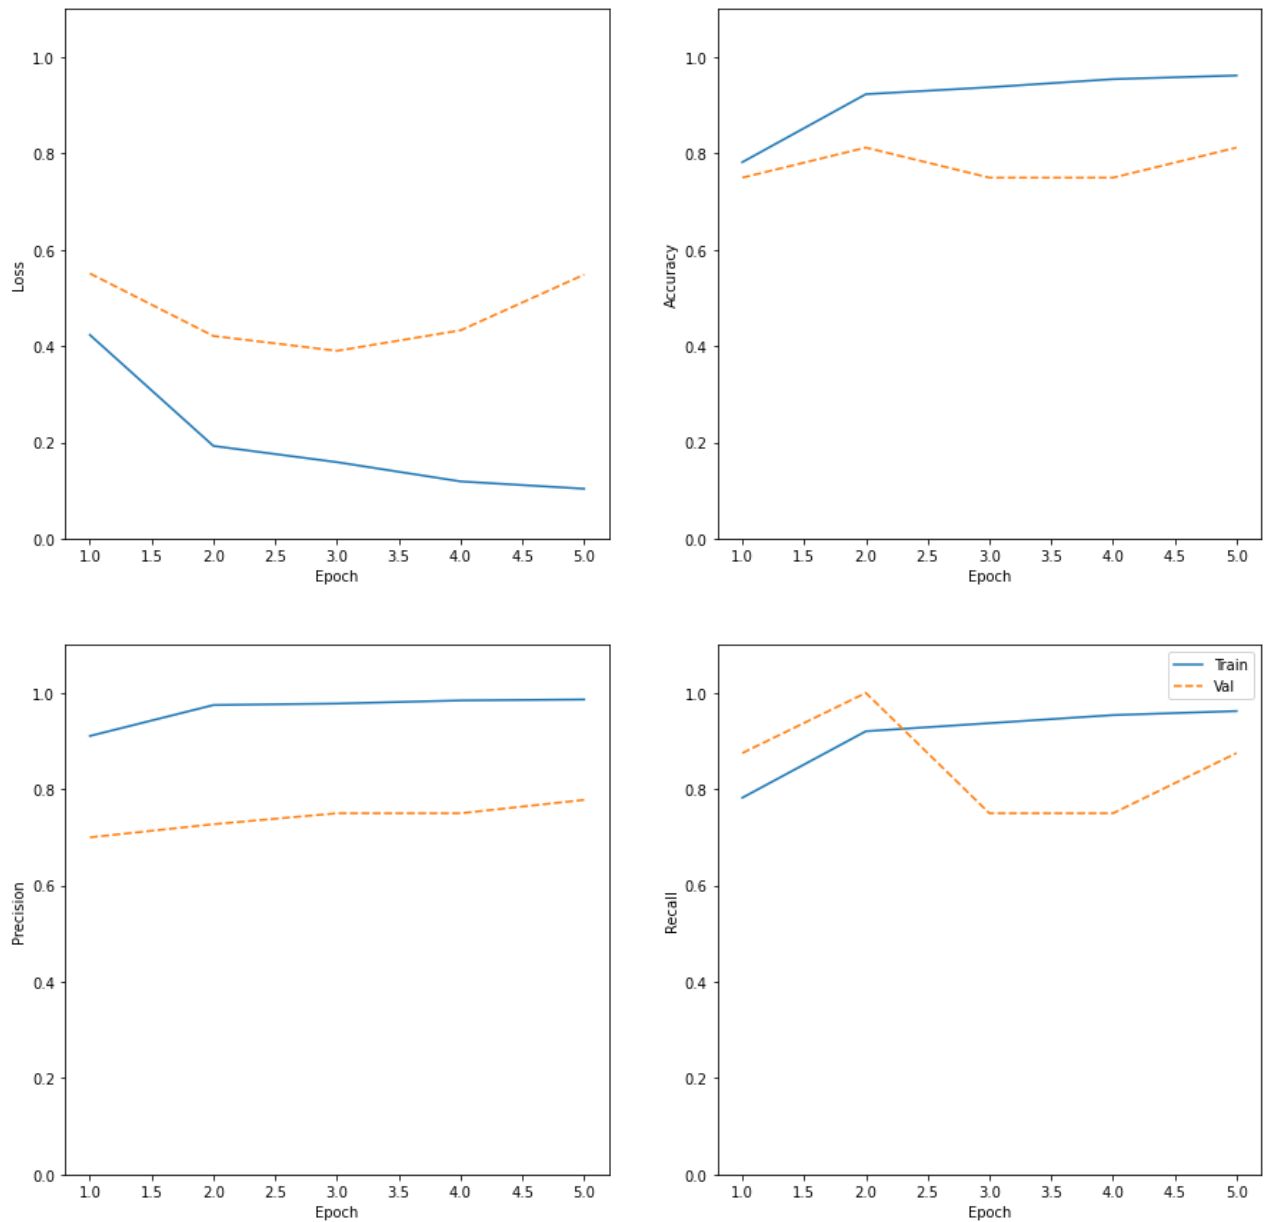
Input V for validation data:

t

```
-----
```

	precision	recall	f1-score	support
normal	0.91	0.52	0.66	234
pneumonia	0.77	0.97	0.86	390
accuracy			0.80	624
macro avg	0.84	0.75	0.76	624
weighted avg	0.82	0.80	0.79	624





CNN model v8

Same model as version 7 with decreased batch size

```
In [60]: def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(16, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))
```

```

model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v8 = make_cnn_model()
cnn_model_v8.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		
conv2d_26 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_26 (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_27 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_27 (MaxPooling2D)	(None, 62, 62, 16)	0
conv2d_28 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_28 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_29 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_29 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten_7 (Flatten)	(None, 12544)	0
dense_18 (Dense)	(None, 256)	3211520
dense_19 (Dense)	(None, 1)	257
=====		
Total params: 3,237,393		
Trainable params: 3,237,393		
Non-trainable params: 0		

In [61]:

```

cnn_model_v8_results = cnn_model_v8.fit(X_train, y_train,
                                         batch_size=64,
                                         epochs=EPOCHS,
                                         callbacks=CALLBACKS,
                                         class_weight=class_weights,
                                         validation_data=(X_val, y_val))

```

Epoch 1/25

```

81/81 [=====] - ETA: 0s - loss: 0.2935 - tp: 3124.0000 - fp: 10
2.0000 - tn: 1247.0000 - fn: 701.0000 - accuracy: 0.8448 - precision: 0.9684 - recall:
0.8167 - auc: 0.9510 - prc: 0.9817
Epoch 1: val_loss did not improve from 0.25460
81/81 [=====] - 82s 982ms/step - loss: 0.2935 - tp: 3124.0000 -
fp: 102.0000 - tn: 1247.0000 - fn: 701.0000 - accuracy: 0.8448 - precision: 0.9684 - rec
all: 0.8167 - auc: 0.9510 - prc: 0.9817 - val_loss: 1.1838 - val_tp: 8.0000 - val_fp: 7.
0000 - val_tn: 1.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.5625 - val_precision: 0.533
3 - val_recall: 1.0000 - val_auc: 0.9219 - val_prc: 0.9264
Epoch 2/25
81/81 [=====] - ETA: 0s - loss: 0.1519 - tp: 3614.0000 - fp: 7
8.0000 - tn: 1263.0000 - fn: 203.0000 - accuracy: 0.9455 - precision: 0.9789 - recall:
0.9468 - auc: 0.9858 - prc: 0.9945
Epoch 2: val_loss did not improve from 0.25460
81/81 [=====] - 79s 971ms/step - loss: 0.1519 - tp: 3614.0000 -
fp: 78.0000 - tn: 1263.0000 - fn: 203.0000 - accuracy: 0.9455 - precision: 0.9789 - reca
ll: 0.9468 - auc: 0.9858 - prc: 0.9945 - val_loss: 0.5065 - val_tp: 8.0000 - val_fp: 3.0
000 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.9375 - val_prc: 0.9265
Epoch 3/25
81/81 [=====] - ETA: 0s - loss: 0.1069 - tp: 3682.0000 - fp: 5
6.0000 - tn: 1285.0000 - fn: 135.0000 - accuracy: 0.9630 - precision: 0.9850 - recall:
0.9646 - auc: 0.9920 - prc: 0.9970
Epoch 3: val_loss did not improve from 0.25460
81/81 [=====] - 79s 973ms/step - loss: 0.1069 - tp: 3682.0000 -
fp: 56.0000 - tn: 1285.0000 - fn: 135.0000 - accuracy: 0.9630 - precision: 0.9850 - reca
ll: 0.9646 - auc: 0.9920 - prc: 0.9970 - val_loss: 0.3370 - val_tp: 7.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - v
al_recall: 0.8750 - val_auc: 0.9531 - val_prc: 0.9493
Epoch 4/25
81/81 [=====] - ETA: 0s - loss: 0.0964 - tp: 3682.0000 - fp: 4
6.0000 - tn: 1295.0000 - fn: 135.0000 - accuracy: 0.9649 - precision: 0.9877 - recall:
0.9646 - auc: 0.9930 - prc: 0.9974
Epoch 4: val_loss did not improve from 0.25460
81/81 [=====] - 80s 991ms/step - loss: 0.0964 - tp: 3682.0000 -
fp: 46.0000 - tn: 1295.0000 - fn: 135.0000 - accuracy: 0.9649 - precision: 0.9877 - reca
ll: 0.9646 - auc: 0.9930 - prc: 0.9974 - val_loss: 0.7860 - val_tp: 8.0000 - val_fp: 6.0
000 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714
- val_recall: 1.0000 - val_auc: 0.9141 - val_prc: 0.9003
Epoch 5/25
81/81 [=====] - ETA: 0s - loss: 0.0853 - tp: 3709.0000 - fp: 3
6.0000 - tn: 1305.0000 - fn: 108.0000 - accuracy: 0.9721 - precision: 0.9904 - recall:
0.9717 - auc: 0.9945 - prc: 0.9979
Epoch 5: val_loss did not improve from 0.25460
81/81 [=====] - 79s 971ms/step - loss: 0.0853 - tp: 3709.0000 -
fp: 36.0000 - tn: 1305.0000 - fn: 108.0000 - accuracy: 0.9721 - precision: 0.9904 - reca
ll: 0.9717 - auc: 0.9945 - prc: 0.9979 - val_loss: 0.3375 - val_tp: 7.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - v
al_recall: 0.8750 - val_auc: 0.9219 - val_prc: 0.9149

```

In [62]:

```
full_report(cnn_model_v8, cnn_model_v8_results)
```

Choose which testing dataset to use for reporting


Input S for training data





Input T for test data

Input V for validation data:

t

	phase_4_project			
	precision	recall	f1-score	support
normal	0.93	0.54	0.68	234
pneumonia	0.78	0.97	0.87	390
accuracy			0.81	624
macro avg	0.85	0.76	0.78	624
weighted avg	0.84	0.81	0.80	624





CNN model v9

Version 8 was the best one yet! Increase in total accuracy to 81%. It does come with a cost of increasing type II errors. Lets increase the batch size to double what we used in version 7

```
In [63]: def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(16, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
```

```
model.add(MaxPooling2D((2,2)))

model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v9 = make_cnn_model()
cnn_model_v9.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
=====		
conv2d_30 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_30 (MaxPoolin g2D)	(None, 127, 127, 16)	0
conv2d_31 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_31 (MaxPoolin g2D)	(None, 62, 62, 16)	0
conv2d_32 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_32 (MaxPoolin g2D)	(None, 30, 30, 32)	0
conv2d_33 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_33 (MaxPoolin g2D)	(None, 14, 14, 64)	0
flatten_8 (Flatten)	(None, 12544)	0
dense_20 (Dense)	(None, 256)	3211520
dense_21 (Dense)	(None, 1)	257
=====		
Total params: 3,237,393		
Trainable params: 3,237,393		
Non-trainable params: 0		

```
In [64]: cnn_model_v9_results = cnn_model_v9.fit(X_train, y_train,
                                                batch_size=256,
                                                epochs=EPOCHS,
                                                callbacks=CALLBACKS,
                                                class_weight=class_weights,
                                                validation_data=(X_val, y_val))
```

```

Epoch 1/25
21/21 [=====] - ETA: 0s - loss: 0.5910 - tp: 1998.0000 - fp: 17
4.0000 - tn: 1175.0000 - fn: 1827.0000 - accuracy: 0.6133 - precision: 0.9199 - recall:
0.5224 - auc: 0.7973 - prc: 0.9064
Epoch 1: val_loss did not improve from 0.25460
21/21 [=====] - 81s 4s/step - loss: 0.5910 - tp: 1998.0000 - f
p: 174.0000 - tn: 1175.0000 - fn: 1827.0000 - accuracy: 0.6133 - precision: 0.9199 - rec
all: 0.5224 - auc: 0.7973 - prc: 0.9064 - val_loss: 0.5301 - val_tp: 3.0000 - val_fp: 0.
0000e+00 - val_tn: 8.0000 - val_fn: 5.0000 - val_accuracy: 0.6875 - val_precision: 1.000
0 - val_recall: 0.3750 - val_auc: 0.9531 - val_prc: 0.9643
Epoch 2/25
21/21 [=====] - ETA: 0s - loss: 0.3104 - tp: 3270.0000 - fp: 14
1.0000 - tn: 1200.0000 - fn: 547.0000 - accuracy: 0.8666 - precision: 0.9587 - recall:
0.8567 - auc: 0.9470 - prc: 0.9804
Epoch 2: val_loss did not improve from 0.25460
21/21 [=====] - 78s 4s/step - loss: 0.3104 - tp: 3270.0000 - f
p: 141.0000 - tn: 1200.0000 - fn: 547.0000 - accuracy: 0.8666 - precision: 0.9587 - reca
ll: 0.8567 - auc: 0.9470 - prc: 0.9804 - val_loss: 0.2955 - val_tp: 7.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - v
al_recall: 0.8750 - val_auc: 0.9844 - val_prc: 0.9853
Epoch 3/25
21/21 [=====] - ETA: 0s - loss: 0.1906 - tp: 3503.0000 - fp: 8
5.0000 - tn: 1256.0000 - fn: 314.0000 - accuracy: 0.9226 - precision: 0.9763 - recall:
0.9177 - auc: 0.9786 - prc: 0.9925
Epoch 3: val_loss did not improve from 0.25460
21/21 [=====] - 78s 4s/step - loss: 0.1906 - tp: 3503.0000 - f
p: 85.0000 - tn: 1256.0000 - fn: 314.0000 - accuracy: 0.9226 - precision: 0.9763 - recal
l: 0.9177 - auc: 0.9786 - prc: 0.9925 - val_loss: 0.3217 - val_tp: 6.0000 - val_fp: 1.00
00 - val_tn: 7.0000 - val_fn: 2.0000 - val_accuracy: 0.8125 - val_precision: 0.8571 - va
l_recall: 0.7500 - val_auc: 0.9531 - val_prc: 0.9493
Epoch 4/25
21/21 [=====] - ETA: 0s - loss: 0.1344 - tp: 3607.0000 - fp: 5
9.0000 - tn: 1282.0000 - fn: 210.0000 - accuracy: 0.9478 - precision: 0.9839 - recall:
0.9450 - auc: 0.9889 - prc: 0.9959
Epoch 4: val_loss did not improve from 0.25460
21/21 [=====] - 79s 4s/step - loss: 0.1344 - tp: 3607.0000 - f
p: 59.0000 - tn: 1282.0000 - fn: 210.0000 - accuracy: 0.9478 - precision: 0.9839 - recal
l: 0.9450 - auc: 0.9889 - prc: 0.9959 - val_loss: 0.6329 - val_tp: 8.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.8125 - val_precision: 0.7273
- val_recall: 1.0000 - val_auc: 0.8594 - val_prc: 0.8686

```

In [65]: `full_report(cnn_model_v9, cnn_model_v9_results)`

Choose which testing dataset to use for reporting






Input S for training data

Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.93	0.38	0.54	234
pneumonia	0.73	0.98	0.84	390
accuracy			0.76	624
macro avg	0.83	0.68	0.69	624
weighted avg	0.80	0.76	0.73	624



CNN model v10

We are going to take version 8 and increase the complexity of our model. We will leave the same number of convolution layers but increase the filters on each layer by double.

In [66]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=OPTIMIZER,
                  metrics=METRICS)

    return model

cnn_model_v10 = make_cnn_model()
cnn_model_v10.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====		
conv2d_34 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_34 (MaxPoolin g2D)	(None, 127, 127, 32)	0
conv2d_35 (Conv2D)	(None, 125, 125, 32)	9248
max_pooling2d_35 (MaxPoolin g2D)	(None, 62, 62, 32)	0
conv2d_36 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d_36 (MaxPoolin g2D)	(None, 30, 30, 64)	0
conv2d_37 (Conv2D)	(None, 28, 28, 128)	73856

max_pooling2d_37 (MaxPoolin g2D)	(None, 14, 14, 128)	0
flatten_9 (Flatten)	(None, 25088)	0
dense_22 (Dense)	(None, 256)	6422784
dense_23 (Dense)	(None, 1)	257

```

=====
Total params: 6,524,961
Trainable params: 6,524,961
Non-trainable params: 0

```

In [67]:

```

cnn_model_v10_results = cnn_model_v10.fit(X_train, y_train,
                                           batch_size=BATCH_SIZE,
                                           epochs=EPOCHS,
                                           callbacks=CALLBACKS,
                                           class_weight=class_weights,
                                           validation_data=(X_val, y_val))

```

Epoch 1/25

```

41/41 [=====] - ETA: 0s - loss: 0.5282 - tp: 2558.0000 - fp: 32
3.0000 - tn: 1026.0000 - fn: 1267.0000 - accuracy: 0.6927 - precision: 0.8879 - recall:
0.6688 - auc: 0.8123 - prc: 0.9193

```

Epoch 1: val_loss did not improve from 0.25460

```

41/41 [=====] - 148s 4s/step - loss: 0.5282 - tp: 2558.0000 - f
p: 323.0000 - tn: 1026.0000 - fn: 1267.0000 - accuracy: 0.6927 - precision: 0.8879 - rec
all: 0.6688 - auc: 0.8123 - prc: 0.9193 - val_loss: 0.3709 - val_tp: 6.0000 - val_fp: 0.
0000e+00 - val_tn: 8.0000 - val_fn: 2.0000 - val_accuracy: 0.8750 - val_precision: 1.000
0 - val_recall: 0.7500 - val_auc: 0.9531 - val_prc: 0.9643

```

Epoch 2/25

```

41/41 [=====] - ETA: 0s - loss: 0.1958 - tp: 3509.0000 - fp: 10
4.0000 - tn: 1237.0000 - fn: 308.0000 - accuracy: 0.9201 - precision: 0.9712 - recall:
0.9193 - auc: 0.9770 - prc: 0.9914

```

Epoch 2: val_loss did not improve from 0.25460

```

41/41 [=====] - 147s 4s/step - loss: 0.1958 - tp: 3509.0000 - f
p: 104.0000 - tn: 1237.0000 - fn: 308.0000 - accuracy: 0.9201 - precision: 0.9712 - reca
ll: 0.9193 - auc: 0.9770 - prc: 0.9914 - val_loss: 0.5579 - val_tp: 8.0000 - val_fp: 4.0
000 - val_tn: 4.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.7500 - val_precision: 0.6667
- val_recall: 1.0000 - val_auc: 0.8906 - val_prc: 0.9014

```

Epoch 3/25

```

41/41 [=====] - ETA: 0s - loss: 0.1398 - tp: 3612.0000 - fp: 7
7.0000 - tn: 1264.0000 - fn: 205.0000 - accuracy: 0.9453 - precision: 0.9791 - recall:
0.9463 - auc: 0.9878 - prc: 0.9955

```

Epoch 3: val_loss did not improve from 0.25460

```

41/41 [=====] - 179s 4s/step - loss: 0.1398 - tp: 3612.0000 - f
p: 77.0000 - tn: 1264.0000 - fn: 205.0000 - accuracy: 0.9453 - precision: 0.9791 - recal
l: 0.9463 - auc: 0.9878 - prc: 0.9955 - val_loss: 0.6484 - val_tp: 7.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 1.0000 - val_accuracy: 0.7500 - val_precision: 0.7000 - va
l_recall: 0.8750 - val_auc: 0.8359 - val_prc: 0.8158

```

In [68]:

```

full_report(cnn_model_v10, cnn_model_v10_results)

```

Choose which testing dataset to use for reporting


Input S for training data


Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.88	0.44	0.59	234
pneumonia	0.74	0.96	0.84	390
accuracy			0.77	624
macro avg	0.81	0.70	0.72	624
weighted avg	0.79	0.77	0.75	624





CNN model v11 - current best overall accuracy

Still not too much improvement off our best model, lets add one more layer of convolution with 32 filters to our model v8.

In [69]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(16, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
```



```

model.add(MaxPooling2D((2,2)))

model.add(Conv2D(64, (3,3), activation=ACTIVATION))
model.add(MaxPooling2D((2,2)))

model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v11 = make_cnn_model()
cnn_model_v11.summary()

```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====		
conv2d_38 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_38 (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_39 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_39 (MaxPooling2D)	(None, 62, 62, 16)	0
conv2d_40 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_40 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_41 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d_41 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_42 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_42 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_10 (Flatten)	(None, 2304)	0
dense_24 (Dense)	(None, 256)	590080
dense_25 (Dense)	(None, 1)	257
=====		
Total params: 625,201		
Trainable params: 625,201		
Non-trainable params: 0		

In [70]:

```
cnn_model_v11_results = cnn_model_v11.fit(X_train, y_train,
                                          batch_size=BATCH_SIZE,
                                          epochs=EPOCHS,
                                          callbacks=CALLBACKS,
                                          class_weight=class_weights,
                                          validation_data=(X_val, y_val))
```

Epoch 1/25

```
41/41 [=====] - ETA: 0s - loss: 0.5210 - tp: 2086.0000 - fp: 15
1.0000 - tn: 1198.0000 - fn: 1739.0000 - accuracy: 0.6347 - precision: 0.9325 - recall:
0.5454 - auc: 0.8434 - prc: 0.9345
```

Epoch 1: val_loss did not improve from 0.25460

```
41/41 [=====] - 85s 2s/step - loss: 0.5210 - tp: 2086.0000 - f
p: 151.0000 - tn: 1198.0000 - fn: 1739.0000 - accuracy: 0.6347 - precision: 0.9325 - rec
all: 0.5454 - auc: 0.8434 - prc: 0.9345 - val_loss: 0.2822 - val_tp: 7.0000 - val_fp: 0.
0000e+00 - val_tn: 8.0000 - val_fn: 1.0000 - val_accuracy: 0.9375 - val_precision: 1.000
0 - val_recall: 0.8750 - val_auc: 1.0000 - val_prc: 1.0000
```

Epoch 2/25

```
41/41 [=====] - ETA: 0s - loss: 0.1963 - tp: 3499.0000 - fp: 9
6.0000 - tn: 1245.0000 - fn: 318.0000 - accuracy: 0.9197 - precision: 0.9733 - recall:
0.9167 - auc: 0.9770 - prc: 0.9921
```

Epoch 2: val_loss did not improve from 0.25460

```
41/41 [=====] - 91s 2s/step - loss: 0.1963 - tp: 3499.0000 - f
p: 96.0000 - tn: 1245.0000 - fn: 318.0000 - accuracy: 0.9197 - precision: 0.9733 - recal
l: 0.9167 - auc: 0.9770 - prc: 0.9921 - val_loss: 0.2719 - val_tp: 8.0000 - val_fp: 0.00
00e+00 - val_tn: 8.0000 - val_fn: 0.0000e+00 - val_accuracy: 1.0000 - val_precision: 1.0
000 - val_recall: 1.0000 - val_auc: 1.0000 - val_prc: 1.0000
```

Epoch 3/25

```
41/41 [=====] - ETA: 0s - loss: 0.1719 - tp: 3557.0000 - fp: 9
0.0000 - tn: 1251.0000 - fn: 260.0000 - accuracy: 0.9321 - precision: 0.9753 - recall:
0.9319 - auc: 0.9822 - prc: 0.9938
```

Epoch 3: val_loss did not improve from 0.25460

```
41/41 [=====] - 86s 2s/step - loss: 0.1719 - tp: 3557.0000 - f
p: 90.0000 - tn: 1251.0000 - fn: 260.0000 - accuracy: 0.9321 - precision: 0.9753 - recal
l: 0.9319 - auc: 0.9822 - prc: 0.9938 - val_loss: 0.3318 - val_tp: 7.0000 - val_fp: 3.00
00 - val_tn: 5.0000 - val_fn: 1.0000 - val_accuracy: 0.7500 - val_precision: 0.7000 - va
l_recall: 0.8750 - val_auc: 0.9531 - val_prc: 0.9643
```

Epoch 4/25

```
41/41 [=====] - ETA: 0s - loss: 0.1148 - tp: 3641.0000 - fp: 4
7.0000 - tn: 1294.0000 - fn: 176.0000 - accuracy: 0.9568 - precision: 0.9873 - recall:
0.9539 - auc: 0.9917 - prc: 0.9972
```

Epoch 4: val_loss did not improve from 0.25460

```
41/41 [=====] - 87s 2s/step - loss: 0.1148 - tp: 3641.0000 - f
p: 47.0000 - tn: 1294.0000 - fn: 176.0000 - accuracy: 0.9568 - precision: 0.9873 - recal
l: 0.9539 - auc: 0.9917 - prc: 0.9972 - val_loss: 0.6745 - val_tp: 8.0000 - val_fp: 6.00
00 - val_tn: 2.0000 - val_fn: 0.0000e+00 - val_accuracy: 0.6250 - val_precision: 0.5714
- val_recall: 1.0000 - val_auc: 0.9062 - val_prc: 0.9333
```

In [71]:

```
full_report(cnn_model_v11, cnn_model_v11_results)
```

Choose which testing dataset to use for reporting

Input S for training data


Input T for test data





Input V for validation data:

t

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

normal	0.93	0.32	0.48	234
pneumonia	0.71	0.98	0.82	390
accuracy			0.74	624
macro avg	0.82	0.65	0.65	624
weighted avg	0.79	0.74	0.70	624





The first time this model was trained it was the best overall. The computer restarted and in order to prepare the necessary deliverables I needed to rerun all the cells. This time the version 11 model performed much worse than originally seen.

```
In [ ]: # cnn_model_v11.save(filepath="./models/")
```

CNN model v12 (11.1)

Version 11 was pretty good, and our early stopping callback was not triggered. Lets run it for more epochs to see if we can get a new best model.

*this statement is no longer valid given the changes when we reran the training for model v11. Thankfully, that model was saved for us to bring back in at the end of the notebook. We also therefore don't need to run this model as it's the same as the last.

```
In [72]: def make_cnn_model():
```

```

model = Sequential()

model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
model.add(MaxPooling2D((2,2)))

model.add(Conv2D(16, (3,3), activation=ACTIVATION))
model.add(MaxPooling2D((2,2)))

model.add(Conv2D(32, (3,3), activation=ACTIVATION))
model.add(MaxPooling2D((2,2)))

model.add(Conv2D(32, (3,3), activation=ACTIVATION))
model.add(MaxPooling2D((2,2)))

model.add(Conv2D(64, (3,3), activation=ACTIVATION))
model.add(MaxPooling2D((2,2)))

model.add(Flatten())

model.add(Dense(256, activation=ACTIVATION))

model.add(Dense(1, activation=OUTPUT))

model.compile(loss=LOSS,
              optimizer=OPTIMIZER,
              metrics=METRICS)

return model

cnn_model_v12 = make_cnn_model()
cnn_model_v12.summary()

```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
=====		
conv2d_43 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_43 (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_44 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_44 (MaxPooling2D)	(None, 62, 62, 16)	0
conv2d_45 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_45 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_46 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d_46 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_47 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_47 (MaxPooling2D)	(None, 6, 6, 64)	0

g2D)

flatten_11 (Flatten)	(None, 2304)	0
dense_26 (Dense)	(None, 256)	590080
dense_27 (Dense)	(None, 1)	257

```

=====
Total params: 625,201
Trainable params: 625,201
Non-trainable params: 0

```

```

In [73]: # cnn_model_v12_results = cnn_model_v12.fit(X_train, y_train,
#                                                  batch_size=BATCH_SIZE,
#                                                  epochs=EPOCHS,
#                                                  callbacks=CALLBACKS,
#                                                  class_weight=class_weights,
#                                                  validation_data=(X_val, y_val))

```

```

In [74]: # full_report(cnn_model_v12, cnn_model_v12_results)

```

CNN model v13 (11.2)

Lets try increasing the value of each layers filter by double.

```

In [75]: def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(32, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(128, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=OPTIMIZER,
                  metrics=METRICS)

```

```
return model
```

```
cnn_model_v13 = make_cnn_model()
cnn_model_v13.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
=====		
conv2d_48 (Conv2D)	(None, 254, 254, 32)	320
max_pooling2d_48 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_49 (Conv2D)	(None, 125, 125, 32)	9248
max_pooling2d_49 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_50 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d_50 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_51 (Conv2D)	(None, 28, 28, 64)	36928
max_pooling2d_51 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_52 (Conv2D)	(None, 12, 12, 128)	73856
max_pooling2d_52 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten_12 (Flatten)	(None, 4608)	0
dense_28 (Dense)	(None, 256)	1179904
dense_29 (Dense)	(None, 1)	257
=====		
Total params: 1,319,009		
Trainable params: 1,319,009		
Non-trainable params: 0		

In [76]:

```
cnn_model_v13_results = cnn_model_v13.fit(X_train, y_train,
                                           batch_size=BATCH_SIZE,
                                           epochs=EPOCHS,
                                           callbacks=CALLBACKS,
                                           class_weight=class_weights,
                                           validation_data=(X_val, y_val))
```

Epoch 1/25

41/41 [=====] - ETA: 0s - loss: 0.4991 - tp: 2045.0000 - fp: 146.0000 - tn: 1203.0000 - fn: 1780.0000 - accuracy: 0.6278 - precision: 0.9334 - recall: 0.5346 - auc: 0.8402 - prc: 0.9335

Epoch 1: val_loss did not improve from 0.25460

```

41/41 [=====] - 165s 4s/step - loss: 0.4991 - tp: 2045.0000 - f
p: 146.0000 - tn: 1203.0000 - fn: 1780.0000 - accuracy: 0.6278 - precision: 0.9334 - rec
all: 0.5346 - auc: 0.8402 - prc: 0.9335 - val_loss: 0.3772 - val_tp: 7.0000 - val_fp: 0.
0000e+00 - val_tn: 8.0000 - val_fn: 1.0000 - val_accuracy: 0.9375 - val_precision: 1.000
0 - val_recall: 0.8750 - val_auc: 0.8906 - val_prc: 0.9396
Epoch 2/25
41/41 [=====] - ETA: 0s - loss: 0.1986 - tp: 3516.0000 - fp: 10
1.0000 - tn: 1240.0000 - fn: 301.0000 - accuracy: 0.9221 - precision: 0.9721 - recall:
0.9211 - auc: 0.9761 - prc: 0.9913
Epoch 2: val_loss did not improve from 0.25460
41/41 [=====] - 171s 4s/step - loss: 0.1986 - tp: 3516.0000 - f
p: 101.0000 - tn: 1240.0000 - fn: 301.0000 - accuracy: 0.9221 - precision: 0.9721 - reca
ll: 0.9211 - auc: 0.9761 - prc: 0.9913 - val_loss: 0.3719 - val_tp: 7.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - v
al_recall: 0.8750 - val_auc: 0.8906 - val_prc: 0.9396
Epoch 3/25
41/41 [=====] - ETA: 0s - loss: 0.1560 - tp: 3579.0000 - fp: 7
7.0000 - tn: 1264.0000 - fn: 238.0000 - accuracy: 0.9389 - precision: 0.9789 - recall:
0.9376 - auc: 0.9851 - prc: 0.9944
Epoch 3: val_loss did not improve from 0.25460
41/41 [=====] - 176s 4s/step - loss: 0.1560 - tp: 3579.0000 - f
p: 77.0000 - tn: 1264.0000 - fn: 238.0000 - accuracy: 0.9389 - precision: 0.9789 - recal
l: 0.9376 - auc: 0.9851 - prc: 0.9944 - val_loss: 0.4429 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.8750 - val_prc: 0.9229
Epoch 4/25
41/41 [=====] - ETA: 0s - loss: 0.1176 - tp: 3632.0000 - fp: 5
6.0000 - tn: 1285.0000 - fn: 185.0000 - accuracy: 0.9533 - precision: 0.9848 - recall:
0.9515 - auc: 0.9909 - prc: 0.9966
Epoch 4: val_loss did not improve from 0.25460
41/41 [=====] - 168s 4s/step - loss: 0.1176 - tp: 3632.0000 - f
p: 56.0000 - tn: 1285.0000 - fn: 185.0000 - accuracy: 0.9533 - precision: 0.9848 - recal
l: 0.9515 - auc: 0.9909 - prc: 0.9966 - val_loss: 0.5132 - val_tp: 7.0000 - val_fp: 4.00
00 - val_tn: 4.0000 - val_fn: 1.0000 - val_accuracy: 0.6875 - val_precision: 0.6364 - va
l_recall: 0.8750 - val_auc: 0.8750 - val_prc: 0.8937

```

In [77]: `full_report(cnn_model_v13, cnn_model_v13_results)`

Choose which testing dataset to use for reporting






Input S for training data

Input T for test data

Input V for validation data:

t

	precision	recall	f1-score	support
normal	0.91	0.41	0.57	234
pneumonia	0.73	0.98	0.84	390
accuracy			0.76	624
macro avg	0.82	0.69	0.70	624
weighted avg	0.80	0.76	0.74	624



CNN model v14 (11.3)

Lets add some dropout to our version 11 model.

In [78]:

```
def make_cnn_model():

    model = Sequential()

    model.add(Conv2D(16, (3,3), activation=ACTIVATION, input_shape=INPUT_SHAPE))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(16, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(32, (3,3), activation=ACTIVATION))
    model.add(MaxPooling2D((2,2)))

    model.add(Conv2D(64, (3,3), activation=ACTIVATION))

    model.add(Flatten())

    model.add(Dense(256, activation=ACTIVATION))
    model.add(Dropout(0.3))

    model.add(Dense(1, activation=OUTPUT))

    model.compile(loss=LOSS,
                  optimizer=OPTIMIZER,
                  metrics=METRICS)

    return model

cnn_model_v14 = make_cnn_model()
cnn_model_v14.summary()
```

Model: "sequential_14"

Layer (type)	Output Shape	Param #
=====		
conv2d_53 (Conv2D)	(None, 254, 254, 16)	160
max_pooling2d_53 (MaxPoolin g2D)	(None, 127, 127, 16)	0
conv2d_54 (Conv2D)	(None, 125, 125, 16)	2320
max_pooling2d_54 (MaxPoolin g2D)	(None, 62, 62, 16)	0
conv2d_55 (Conv2D)	(None, 60, 60, 32)	4640
max_pooling2d_55 (MaxPoolin g2D)	(None, 30, 30, 32)	0

conv2d_56 (Conv2D)	(None, 28, 28, 32)	9248
max_pooling2d_56 (MaxPoolin g2D)	(None, 14, 14, 32)	0
conv2d_57 (Conv2D)	(None, 12, 12, 64)	18496
flatten_13 (Flatten)	(None, 9216)	0
dense_30 (Dense)	(None, 256)	2359552
dropout_3 (Dropout)	(None, 256)	0
dense_31 (Dense)	(None, 1)	257

```

=====
Total params: 2,394,673
Trainable params: 2,394,673
Non-trainable params: 0

```

In [79]:

```

cnn_model_v14_results = cnn_model_v14.fit(X_train, y_train,
                                           batch_size=BATCH_SIZE,
                                           epochs=EPOCHS,
                                           callbacks=CALLBACKS,
                                           class_weight=class_weights,
                                           validation_data=(X_val, y_val))

```

Epoch 1/25

```

41/41 [=====] - ETA: 0s - loss: 0.3562 - tp: 3104.0000 - fp: 19
1.0000 - tn: 1158.0000 - fn: 721.0000 - accuracy: 0.8237 - precision: 0.9420 - recall:
0.8115 - auc: 0.9215 - prc: 0.9698

```

Epoch 1: val_loss did not improve from 0.25460

```

41/41 [=====] - 90s 2s/step - loss: 0.3562 - tp: 3104.0000 - f
p: 191.0000 - tn: 1158.0000 - fn: 721.0000 - accuracy: 0.8237 - precision: 0.9420 - reca
ll: 0.8115 - auc: 0.9215 - prc: 0.9698 - val_loss: 0.4167 - val_tp: 7.0000 - val_fp: 1.0
000 - val_tn: 7.0000 - val_fn: 1.0000 - val_accuracy: 0.8750 - val_precision: 0.8750 - v
al_recall: 0.8750 - val_auc: 0.8750 - val_prc: 0.9085

```

Epoch 2/25

```

41/41 [=====] - ETA: 0s - loss: 0.1889 - tp: 3505.0000 - fp: 8
3.0000 - tn: 1258.0000 - fn: 312.0000 - accuracy: 0.9234 - precision: 0.9769 - recall:
0.9183 - auc: 0.9780 - prc: 0.9915

```

Epoch 2: val_loss did not improve from 0.25460

```

41/41 [=====] - 86s 2s/step - loss: 0.1889 - tp: 3505.0000 - f
p: 83.0000 - tn: 1258.0000 - fn: 312.0000 - accuracy: 0.9234 - precision: 0.9769 - recal
l: 0.9183 - auc: 0.9780 - prc: 0.9915 - val_loss: 0.4857 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.8438 - val_prc: 0.8816

```

Epoch 3/25

```

41/41 [=====] - ETA: 0s - loss: 0.1337 - tp: 3612.0000 - fp: 6
5.0000 - tn: 1276.0000 - fn: 205.0000 - accuracy: 0.9477 - precision: 0.9823 - recall:
0.9463 - auc: 0.9884 - prc: 0.9957

```

Epoch 3: val_loss did not improve from 0.25460

```


41/41 [=====] - 86s 2s/step - loss: 0.1337 - tp: 3612.0000 - f
p: 65.0000 - tn: 1276.0000 - fn: 205.0000 - accuracy: 0.9477 - precision: 0.9823 - recal
l: 0.9463 - auc: 0.9884 - prc: 0.9957 - val_loss: 0.5309 - val_tp: 7.0000 - val_fp: 2.00
00 - val_tn: 6.0000 - val_fn: 1.0000 - val_accuracy: 0.8125 - val_precision: 0.7778 - va
l_recall: 0.8750 - val_auc: 0.8281 - val_prc: 0.8588

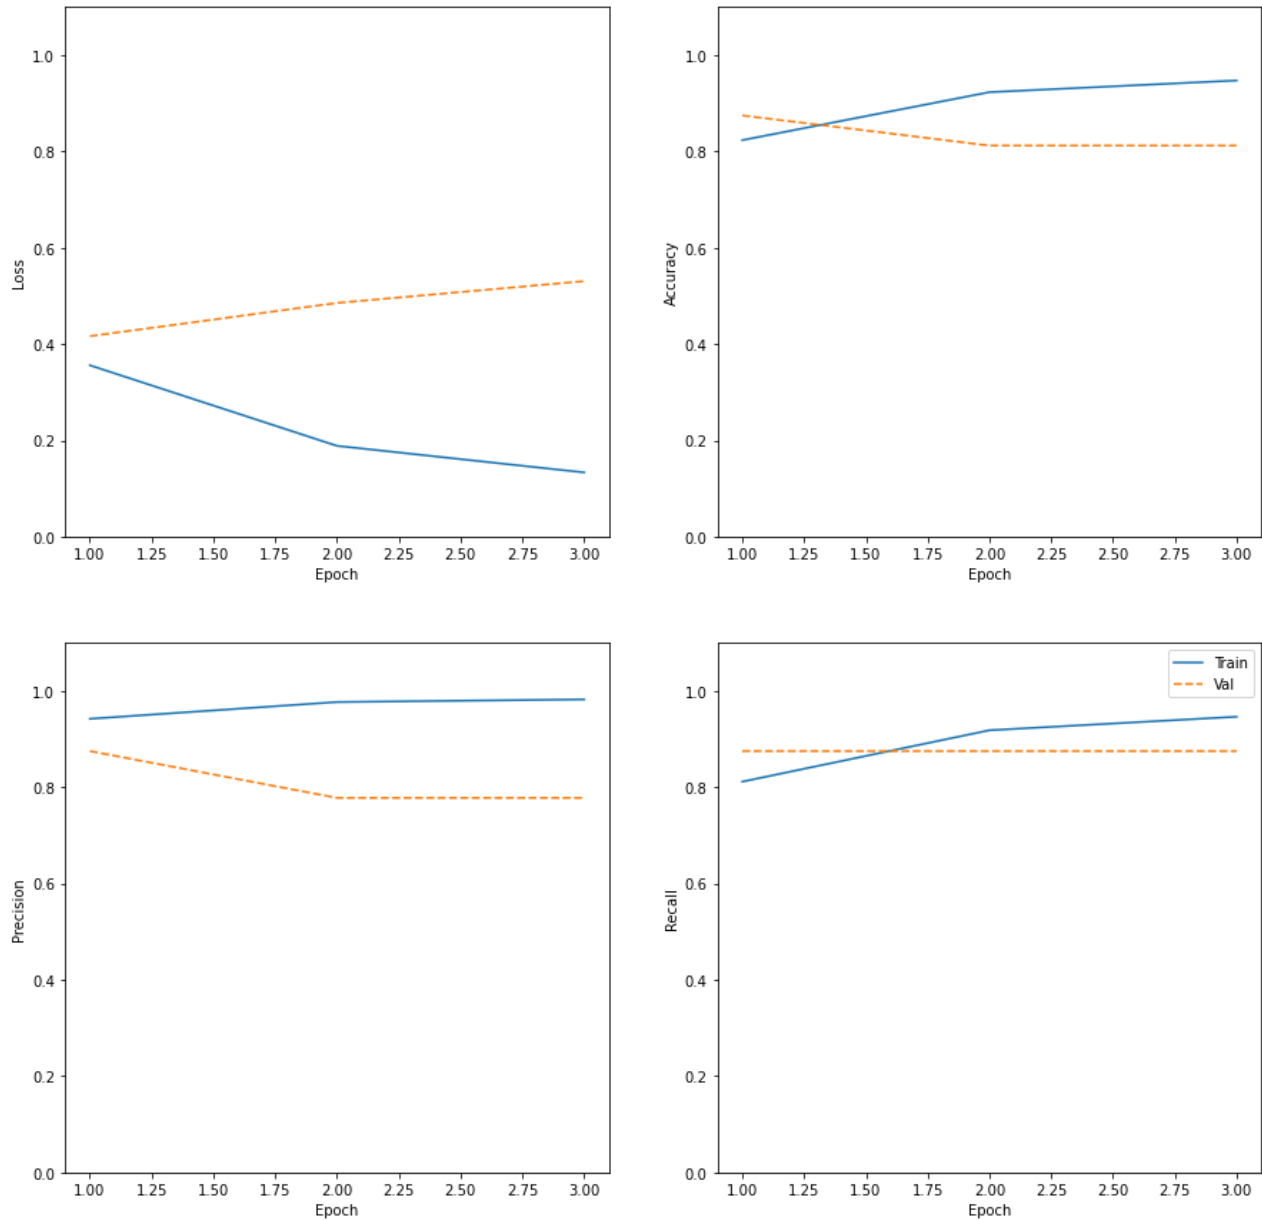
```

```
In [80]: full_report(cnn_model_v14, cnn_model_v14_results)
```

Choose which testing dataset to use for reporting
Input S for training data
Input T for test data
Input V for validation data:
t

	precision	recall	f1-score	support
normal	0.88	0.52	0.65	234
pneumonia	0.77	0.96	0.85	390
accuracy			0.79	624
macro avg	0.82	0.74	0.75	624
weighted avg	0.81	0.79	0.78	624





Final model


```
In [81]: final_model = keras.models.load_model('./models/')
```

```
In [82]: full_report(final_model)
```

Choose which testing dataset to use for reporting
Input S for training data
Input T for test data
Input V for validation data:
t

	precision	recall	f1-score	support
normal	0.85	0.62	0.72	234
pneumonia	0.80	0.94	0.86	390

accuracy			0.82	624
macro avg	0.83	0.78	0.79	624
weighted avg	0.82	0.82	0.81	624



In []: